

Predictive Hard Drive Failure Modeling for SRE Operations

A Cost-Sensitive Machine Learning Framework Bridging Predictive
Maintenance and Site Reliability Engineering

A CAPSTONE PROJECT REPORT PRESENTED
BY
FABIAN HIERNAUX
TO
Harvard Online (HarvardX / EDX)

IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE
**HARVARDX PROFESSIONAL CERTIFICATE
IN DATA SCIENCE**

HARVARD ONLINE (HARVARDX)
CAMBRIDGE, MASSACHUSETTS
September 6, 2026

Executive Summary

I. Context & Operational Challenge

Modern cloud storage infrastructure relies on high-density hard disk drive (HDD) deployments to manage exabyte-scale data. In a Site Reliability Engineering (SRE) context, unannounced drive failures induce operational toil, saturate network bandwidth during ungraceful RAID rebuilds, and risk Service Level Objective (SLO) violations. This capstone project addresses this operational challenge by designing an end-to-end, cost-sensitive machine learning framework using Backblaze Q1 2024 quarterly telemetry (>13 million filtered daily observations) to forecast drive failures prior to catastrophic breakdown. Beyond the predictive modeling objective itself, this project also serves as a practical proof-of-concept for the adoption of Site Reliability Engineering (SRE) principles within infrastructure operations. By demonstrating how large-scale telemetry can be transformed into actionable maintenance decisions, the study evaluates the technical feasibility, operational relevance, and potential business value of predictive maintenance as an enabler of proactive reliability management.

II. Predictive Strategy & System Design

To balance early warning capabilities with high signal confidence, the target failure horizon was determined to be $H = 14$ **days** based on operational considerations. This window provides automated distributed orchestrators (e.g., Ceph, Google Colossus) adequate lead time to gracefully drain data blocks off failing drives while maintaining an optimal signal-to-noise ratio.

To capture non-linear drive degradation dynamics, static point-in-time S.M.A.R.T. metrics were augmented with 7-day dynamic rolling deltas (Δ_7) across critical indicators:

- **Reallocated Sectors Count** (`smart_5_delta7`)
- **Reported Uncorrectable Errors** (`smart_187_delta7`)
- **Current Pending Sector Count** (`smart_197_delta7`)

Additionally, the entire ingestion, feature engineering, and training pipeline was built under strict “Design by Constraint” resource limits (6 GB RAM, 4 vCPUs) to demonstrate the feasibility of lightweight, cost-effective AIOps pipelines.

III. Key Empirical Results

The pipeline was evaluated on an out-of-time temporal test dataset comprising over 4.63 million daily observations with extreme class imbalance ($\approx 0.04\%$ daily failure prevalence):

- **Discriminative Power:** The calibrated XGBoost classifier achieved an **AUC-ROC of 0.8364**, outperforming the baseline Logistic Regression model (0.8176).
- **Imbalanced Precision:** The model reached a **PR-AUC of 0.0516**, representing a **$\sim 129\times$ precision lift** over random guessing on an extreme baseline.
- **Economic Cost Optimization:** Under an enterprise SRE cost matrix penalizing unhandled failures ($C_{FN} = \$500$), decision threshold optimization demonstrated that proactive failure detection drastically reduces emergency on-call toil, network rebuild overhead, and risk of data loss.

IV. Strategic Operational Vision

By converting raw, noisy hardware telemetry into high-confidence proactive alerts, this framework establishes a direct bridge between predictive Data Science and Site Reliability Engineering. The results prove that automated data evacuation (*drain mode*) can be seamlessly integrated into cloud infrastructure pipelines to preserve error budgets and sustain system resilience at enterprise scale.

Acknowledgements

This capstone project represents the culmination of an intensive, self-directed learning journey completed through the edX platform. Operating as an independent learner required deep self-reliance, though I am grateful to the platform instructors and the forum community members whose asynchronous discussions provided valuable peer perspectives along the way.

Pursuing this advanced curriculum via distance learning while balancing active professional responsibilities (combining freelance commitments with business management) alongside family life at a more mature stage of my career was a significant undertaking. Navigating these overlapping demands, late-night study sessions, and work schedules proved to be a grueling balancing act, but ultimately an immensely rewarding personal and intellectual challenge.

I extend my appreciation to **Backblaze** for providing public access to their enterprise-grade telemetry datasets, as well as to the open-source Data Science and Site Reliability Engineering (SRE) communities whose tools and operational frameworks made this work possible.

Finally, my deepest and most heartfelt gratitude goes to my wife and my four children. Completing this capstone required sacrificing countless hours, weekends, and late nights ; often drawn directly from family time and holidays. I am profoundly grateful for their unwavering patience, resilience, and support throughout this journey. This achievement is as much theirs as it is mine.

Contents

Executive Summary	1
I. Context & Operational Challenge	1
II. Predictive Strategy & System Design	1
III. Key Empirical Results	1
IV. Strategic Operational Vision	1
Acknowledgements	2
Glossary & List of Abbreviations	7
1 Introduction	9
1.1 Research Question	9
1.2 Operational and Scientific Context: SRE, Reliability and Proactive Maintenance	10
1.3 Project Objectives & Dataset Overview	11
1.4 Design by Constraint: VM Computational Strategy	11
2 Review of Scientific and Industrial Literature	13
2.1 Empirical Reference Studies: From Google’s Analysis to NetApp’s Failures	13
2.1.1 The Relative Influence of the Physical Environment and Usage	13
2.1.2 Critical S.M.A.R.T. Health Metrics	13
2.2 The challenge of silent corruption and disk typologies	14
2.2.1 Drive types and asymmetric behaviour in response to anomalies	14
2.2.2 Physical localisation and risks of dual failure	15
2.3 Mathematical and methodological theoretical framework in data science	15
2.3.1 Handling Extreme Class Imbalance (Imbalanced Data)	15
2.3.2 Problem formulation: Time-window classification vs. survival analysis	15
2.3.3 Assessing the importance of variables and explainability (XAI)	15
2.4 Systemic convergence: from raw hardware telemetry to sre observability	16
2.4.1 The epistemological shift: from traditional monitoring to proactive observability	16
2.4.2 Operational impact: eliminating Toil and protecting SLOs	16
3 Methods and Analysis	18
3.1 Methodological Framework and Experimental Design	18
3.2 Data Ingestion, Cleaning, Label Alignment and Memory Management Strategy	18
3.2.1 Ingestion and Standardization of Backblaze Telemetry Data	18
3.2.2 From Raw Telemetry to Analytical Readiness	21

4	Exploratory Data Analysis (EDA) of Backblaze Telemetry	22
4.1	Global Failure Patterns and Feature Behaviour	22
4.2	Extended Exploratory Analysis: Correlation, Temporal Dynamics, and Capacity Effects . . .	28
4.2.1	Class Imbalance and Model-Level Vulnerabilities	34
4.2.2	SMART Health Distributions and Collinearity Analysis	35
4.2.3	Temporal Dynamics and Capacity Risk Profiles	37
4.2.4	Predictive Window and Feature Engineering Design	38
5	Machine Learning and Predictive Modelling in an SRE Context	39
5.1	Construction of the Predictive Failure Target ($H = 14$ Days)	39
5.1.1	Strict Temporal Train/Test Splitting Protocol	45
5.1.2	Feature Engineering: The 7-Day SMART Delta (Δ_7)	45
5.2	Model Evaluation: Reliability Metrics, ROC/PR Curves and Temporal Generalisation	46
5.2.1	Discriminative Power (ROC Curves Analysis)	48
5.2.2	Precision-Recall Dynamics in Extreme Class Imbalance	49
5.2.3	Performance Summary	50
5.3	SRE Cost-Benefit Analysis: Operational Economics and Threshold Optimisation	51
5.3.1	Cost Matrix Formulation and Probability Calibration	57
5.3.2	Operational Sensitivity Analysis across Infrastructure Scenarios	58
5.3.3	Economic Results and Financial Impact	59
5.3.4	Strategic Interpretation of Economic Results	59
5.4	Model Interpretability: XGBoost Feature Importance and SHAP-Based Explainability	59
5.4.1	Global Feature Importance and SHAP Framework	63
5.4.2	Domain-Specific Interpretation of Physical Degradation Drivers	64
6	Results and Discussion	65
6.1	Performance Evaluation and Comparative Analysis	65
6.2	Error Analysis and Feature Importance	65
6.3	Operational Trade-offs and Limitations	66
6.4	Feature Importance and SRE Health Indicators	66
6.4.1	Primary Degradation Triggers (<code>smart_197_raw</code> & <code>smart_197_delta7</code>)	66
6.4.2	Physical Remapping and Surface Retirement (<code>smart_5_raw</code> & <code>smart_5_delta7</code>) . . .	66
6.4.3	Immediate Severity Escalators (<code>smart_187_raw</code>)	66
6.4.4	Observability Optimization and Noise Reduction	67
6.4.5	Strategic Synthesis for AIOps Integration	67
6.5	SRE Operational Translation & Business Impact	67
6.5.1	Automated Remediation Orchestration	68
6.5.2	Quantifiable Business Impact and ROI	68

7 Conclusion and Perspectives	70
7.1 Summary of Findings	70
7.2 Operational & Strategic Impact for SRE	70
7.3 Limitations	70
7.4 Future Work & Next Steps	71
References	73
Appendices	74
Appendix A: Overview of the S.M.A.R.T. Protocol and Attribute Taxonomy	74
Physical and Sector Error Metrics	74
Environmental and Mechanical Metrics	74
Interface and System Integration	74

List of Figures

1	Class Imbalance between Operational Logs and 14-Day Failure Horizon	34
2	Daily Failure Rate Heterogeneity across Hard Drive Models	34
3	Log-Transformed SMART Feature Distributions: Failing vs Healthy Drives	35
4	Spearman Correlation Matrix among Key SMART Telemetry Features	36
5	Daily Failure Rate Dynamics and LOESS Trend Line across Q1 2024	37
6	Daily Failure Rate Comparison across Operational Storage Capacities	38
7	ROC Curves Comparison on Temporal Test Set (March 2024)	49
8	XGBoost Precision-Recall Curve under 0.04% Class Imbalance	50
9	SRE Total Maintenance Cost vs. Calibrated Probability Thresholds for Standard (100FN)andEnterpriseSLA(500 FN) Infrastructure Scenarios	58
10	Global Model Interpretability Comparing Native XGBoost Feature Gain (%) with Absolute Mean SHAP Value Contributions	63
11	Automated SRE Remediation Architecture driven by Machine Learning Telemetry.	68

List of Tables

1	Raw Target Variable Distribution (25.1M Daily Logs)	27
2	Top 10 Most Common Hard Drive Models (Raw Telemetry)	27
3	Selected Top 5 Fleet Composition (Processed Dataset)	27
4	Transformed Target Horizon Distribution (Filtered Fleet)	27
5	Model Performance Summary on Out-of-Time Test Set (March 2024)	50
6	SRE Economic Cost Optimization Summary on Out-of-Time Test Set	59
7	Model Performance Summary	65

Glossary & List of Abbreviations

Abbreviation / Term	Definition
AIOps	Artificial Intelligence for IT Operations: The application of machine learning to automate and enhance IT operations and observability workflows.
AUC-ROC	Area Under the Receiver Operating Characteristic Curve: A performance metric measuring a classifier's ability to distinguish between classes across all thresholds.
Backblaze	A cloud storage and data backup provider that publicly releases daily open-source S.M.A.R.T. hard drive failure telemetry datasets.
DevOps	Development and Operations: A cultural and technical movement emphasizing shared responsibility, continuous integration/continuous delivery (CI/CD), and automation to bridge the traditional gap between software development and IT operations teams.
Drain Mode	An operational procedure in distributed storage systems (e.g., Ceph, Google Colossus) that safely migrates data blocks off a drive prior to hardware removal.
HDD / SSD	Hard Disk Drive / Solid State Drive: Non-volatile digital storage technologies; mechanical/magnetic and semiconductor-based, respectively.
PR-AUC	Precision-Recall Area Under the Curve: A robust performance metric specifically suited for evaluating machine learning models on highly imbalanced datasets.
RAID	Redundant Array of Independent Disks: A data storage virtualization technology that combines multiple physical drive components into one logical unit for redundancy and performance.
RUL	Remaining Useful Life: The estimated time (in days or hours) a component is expected to function reliably before experiencing a fatal breakdown.
SLO / SLA	Service Level Objective / Service Level Agreement: SLO defines the target reliability level for a service (e.g., 99.99% availability), while SLA is the commitment to external users including penalties.
SMART / S.M.A.R.T.	Self-Monitoring, Analysis, and Reporting Technology: A monitoring system included in computer hard disk drives that detects and reports indicators of drive reliability.

Abbreviation / Term	Definition
SRE	Site Reliability Engineering: A discipline originating at Google (often summarized as <i>class SRE implements DevOps</i>) that applies software engineering practices to infrastructure and operational problems. It leverages Service Level Objectives (SLOs) and Error Budgets to balance system stability with deployment velocity while actively eliminating toil.
TCN	Temporal Convolutional Network: A type of neural network architecture designed for sequence modeling tasks with causal and dilated convolutions.
Toil	In SRE, manual, repetitive, automatable operational work directly tied to running a service that scales linearly with service growth.
XGBoost	eXtreme Gradient Boosting: An optimized distributed gradient boosting library designed to be highly efficient, flexible, and portable.

1 Introduction

In my professional environment, our organisation is currently engaged in a strategic transition toward Site Reliability Engineering (SRE) practices, with a particular focus on leveraging weak signals and operational telemetry to enable predictive maintenance across critical infrastructure components. This shift reflects a broader ambition: evolving from traditional reactive operations toward a data-driven reliability model in which operational decisions are informed by quantitative evidence, temporal degradation patterns, and machine-learning-based forecasting.

As our internal storage systems do not yet expose publicly accessible historical telemetry suitable for academic experimentation, the HarvardX Data Science Professional Certificate capstone offered an ideal opportunity to explore this challenge using a high-quality, openly available dataset. For this purpose, I selected the Backblaze Hard Drive Stats dataset, one of the most widely used public sources of large-scale S.M.A.R.T. telemetry. Backblaze is a cloud backup provider operating tens of thousands of consumer-grade hard drives, and since 2013 the company has published daily operational snapshots (including failures, model information, and detailed S.M.A.R.T. metrics) to support research in reliability engineering, predictive maintenance, and large-scale storage behaviour analysis.

The choice of topic (predicting imminent hard drive failures within a 14-to-30-day horizon) therefore emerged naturally from the intersection between the academic objectives of the HarvardX programme and the operational challenges currently under discussion within my organisation. This report thus serves a dual purpose: it satisfies the capstone requirements of the HarvardX Data Science Professional Certificate, and it contributes to our ongoing reflection on how data science, observability, and machine learning can support the adoption of modern SRE practices and enable proactive infrastructure reliability. Beyond its academic objectives, the project functions as a feasibility study for predictive maintenance within an SRE transformation context. The intention is not merely to assess forecasting performance, but also to evaluate whether telemetry-driven risk prediction can realistically support operational decision-making, reduce reactive interventions, and contribute to the broader shift toward proactive reliability engineering.

1.1 Research Question

“Can we accurately predict imminent hard drive failures within a 7-to-30-day window using S.M.A.R.T. (Self-Monitoring, Analysis, and Reporting Technology) telemetry, and which specific health indicators are the most critical predictors for proactive hardware replacement?”

To monitor drive health, storage infrastructure relies on S.M.A.R.T. (Self-Monitoring, Analysis and Reporting Technology), an ATA/ATAPI standard providing internal indicators to anticipate mechanical and electronic failures. While the overall framework captures a broad spectrum of metrics (ranging from environmental variables to mechanical fatigue) empirical reliability studies (Pinheiro et al., Google; Schroeder & Gibson, CMU/NetApp) demonstrate that failure prediction is predominantly driven by physical surface degradation and unrecoverable read errors. Consequently, this study filters raw telemetry to focus on five critical indicators: SMART 5 (Reallocated Sectors Count), SMART 187 (Reported Uncorrectable Errors), SMART 196 (Reallocation Event Count), SMART 197 (Current Pending Sector Count), and SMART 198 (Offline Uncorrectable Sectors). These metrics reflect physical sector remapping and hardware ECC limitations, providing the highest predictive signal for impending drive failure.

For a comprehensive breakdown of the S.M.A.R.T. attribute taxonomy and physical degradation mechanisms, see Appendix A.

Formally, let X_{it} represent the vector of S.M.A.R.T. telemetry features for drive i at day t . We formulate hard drive failure prediction as a binary classification problem over a fixed temporal horizon Δt or $H \in \{7, 30\}$ days:

$$Y_{it} = \begin{cases} 1 & \text{if drive } i \text{ fails within } [t + 1, t + \Delta t] \\ 0 & \text{otherwise} \end{cases}$$

While the broader literature explores failure horizons ranging from 7 to 30 days, this study focuses on a 14-day prediction window ($H = 14$), consistent with the target label (target_14d) used throughout the pipeline. In an SRE operational context, a 14-day window strikes an optimal balance: Defining a 7-to-30-day proactive window addresses critical infrastructure requirements, it provides distributed storage orchestrators (e.g., Google’s Colossus) and SRE teams sufficient lead time to transparently drain data blocks from degrading drives to healthy nodes without impacting user-facing Service Level Objectives (SLOs) or incurring emergency maintenance toil (Pineiro et al., 2007; Beyer et al., 2016).

1.2 Operational and Scientific Context: SRE, Reliability and Proactive Maintenance

In today’s cloud infrastructure, hardware failure is no longer an isolated event but a daily statistical certainty. Distributed storage platforms rely on clusters of hundreds of thousands of physical hard disks, which fail continuously (Pineiro et al., 2007; Bairavasundaram et al., 2008). Historically, storage management relied on reactive approaches: replacing drives only after they had actually failed. This traditional model results in a growing volume of maintenance work, emergency interventions and disruptions that are incompatible with the availability requirements of modern systems (Beyer et al., 2016).

Site Reliability Engineering (SRE) offers a methodological shift by applying the principles of software engineering to operational challenges. Within this framework, proactive maintenance becomes a key lever: anticipating a disk failure before it occurs suddenly makes it possible to transform a hardware emergency into a planned operation, integrated into standard maintenance windows.

While the theoretical literature considers extended prediction windows ranging from 7 to 30 days (Pineiro et al., 2007; Beyer et al., 2016), this study adopts a 14-day prediction window ($H = 14$), which balances actionable operational lead time with sufficient positive-class density given the extreme prevalence of failures. This choice is driven by two core operational and technical factors:

- **Signal-to-Noise Ratio and Signal Density:** Physical hard drive failure mechanisms—particularly those reflected in critical S.M.A.R.T. attributes such as pending sector counts (SMART 197) and reallocated sectors (SMART 5)—often exhibit highly localized, non-linear acceleration in the days immediately preceding complete breakdown. Extending the target window beyond 14 days introduces temporal noise and diffuses the predictive signal, leading to a higher rate of false positives. A 14-day lead time captures a sufficiently early period of degradation, significantly maximizing the class separability for tree-based models (Random Forest and XGBoost).
- **Computational Efficiency and Memory Constraints:** From an infrastructure perspective, computing rolling temporal deltas (Δ_7) and lag features across a large single-quarter telemetry dataset already imposes significant RAM overhead during `shift()` and down-sampling operations in R. The 7-day rolling deltas (Δ_7), a feature look-back window distinct from the 14-day target horizon (H), were retained as the shortest window still capturing meaningful degradation acceleration, while keeping memory consumption within the constrained 6 GB RAM profile of the virtual machine. This point regarding the virtual machine’s limitations is explained below in the chapter ‘Design by Constraint: VM Computational Strategy’.

From an SRE operational standpoint, a 14-day window remains fully actionable. It provides automated orchestrators (such as Google’s Colossus) and infrastructure teams sufficient lead time to transparently migrate data blocks and re-balance storage nodes without violating availability SLOs or latency budgets (Thurgood et al., 2018). It also enables the automated creation of maintenance tickets, optimises hardware procurement logistics, and drastically reduces MTTR.

From an operational perspective, the trade-off is asymmetric:

- a **false negative** (unpredicted failure) triggers an immediate RAID rebuild, saturating internal bandwidth and degrading application performance;

- a **false positive** (premature replacement) incurs only a marginal hardware cost.

This asymmetry justifies the use of models capable of detecting weak, non-linear signals in S.M.A.R.T. telemetry, where static thresholds fail. Indeed, Google’s seminal studies show that 56 per cent of failing drives exhibit no anomalies in critical metrics prior to failure (Pinheiro et al., 2007), making the use of temporal and non-linear models essential.

From a modern observability perspective, S.M.A.R.T. data should no longer be viewed as isolated alerts, but as continuous streams of high-cardinality telemetry. Machine learning then acts as an AIOps layer: it correlates temporal trends, identifies subtle deviations, and produces an actionable probability of failure.

In summary, the proactive prediction of hard disk failures lies at the intersection of three dynamics: 1. hardware reliability observed at scale, 2. the operational requirements of distributed systems, 3. and the evolution of SRE practices towards intelligent observability.

The following chapter explores these foundations in greater depth by presenting a structured literature review covering major empirical studies, critical S.M.A.R.T. metrics, silent corruption phenomena, and the theoretical frameworks used in data science for predicting hardware failures.

1.3 Project Objectives & Dataset Overview

Empirical studies from Google (Pinheiro et al., 2007) and NetApp (Bairavasundaram et al., 2008) established that individual S.M.A.R.T. attributes (such as Scan Errors, Reallocation Counts, Offline Reallocations, and Probational Counts) significantly increase 60-day failure probability (up to 39x). However, Pinheiro et al. (2007) also demonstrated that 56% of failed drives exhibited zero warnings across all critical S.M.A.R.T. parameters prior to failure, highlighting the limits of static rules and the need for non-linear machine learning models capable of capturing subtle temporal degradation trends.

To address this challenge, this project leverages real-world telemetry provided by Backblaze, containing daily operational snapshots and S.M.A.R.T. attributes collected from large-scale production storage environments. The objective is to build, evaluate, and interpret supervised machine learning models capable of predicting whether a hard drive will fail within the next 14 days.

To achieve this, the project establishes an end-to-end data pipeline in R:

- Automated data ingestion and temporal filtering of Backblaze quarterly telemetry.
- Exploratory Data Analysis (EDA) focused on extreme class imbalance ($< 1\%$ failure rate), missingness, and failure trends across drive models.
- Feature engineering, including target label propagation (will_fail within H or Δt days) and rolling temporal deltas to capture non-linear telemetry drifts.
- Class imbalance mitigation using strategic down-sampling techniques.
- Comparative modeling evaluating a baseline Logistic Regression model and an advanced XGBoost classifier.
- Model evaluation & Explainability (XAI) using SRE-relevant metrics (Recall, Precision, PR-AUC) and feature importance analysis (e.g., SHAP values) to identify top failure predictors.

1.4 Design by Constraint: VM Computational Strategy

To simulate real-world resource constraints and adhere to strict technical limits, all data processing, feature engineering, and model training were conducted within a restricted Virtual Machine environment equipped with 6 GB of RAM and 4 CPUs.

Operating under restricted memory required an intentional “design by constraint” engineering strategy:

- **Data Optimization:** Ingestion utilizing high-performance I/O (`data.table::fread`), selective column loading, and optimized data types (e.g., integer representation for raw counters).

- **Memory Management:** Explicit garbage collection (`gc()`) invoked after resource-intensive transformations to prevent out-of-memory (OOM) faults.
- **Smart Sampling:** Strategic temporal slicing and structured down-sampling of healthy drive telemetry to ensure low computational overhead without sacrificing statistical rigor.

Having defined the research question, the SRE operational objectives and the general methodological framework of the project, it is necessary to ground this work in existing scientific and industrial literature.

2 Review of Scientific and Industrial Literature

This chapter presents the key research into the reliability of hard disk drives in production, the limitations of S.M.A.R.T. metrics, and the theoretical frameworks used to predict hardware failures.

It draws on two complementary datasets: 1. large-scale empirical analyses conducted by Google (Pinheiro et al., 2007) and NetApp (Bairavasundaram et al., 2008), 2. and conceptual frameworks derived from reliability engineering, data science and SRE practices.

The aim is to situate this work within the state of the art, to identify the most relevant physical signals, and to understand the limitations of traditional approaches in order to justify the use of non-linear and temporal models.

2.1 Empirical Reference Studies: From Google’s Analysis to NetApp’s Failures

Modelling failure prediction draws on a body of first-rate, large-scale empirical studies. The historical analysis published by Google (Pinheiro, Weber & Barroso, 2007), covering a population of over 100,000 consumer-grade hard disk drives (SATA and PATA, ranging from 5,400 to 7,200 rpm, and with capacities from 80 to 400 GB), forms the cornerstone of research in this field. Followed in 2008 by the study by NetApp and the University of Wisconsin–Madison (Bairavasundaram et al., 2008), which analysed more than 1.53 million storage units over a period of 41 months; these two contributions have made it possible to rigorously quantify the physical behaviour of hard disk drives in real-world production environments within professionally managed data centres.

2.1.1 The Relative Influence of the Physical Environment and Usage

One of the most surprising findings of the Google study is the debunking of the myth that temperature is the predominant factor in hardware failure. Whilst historical manufacturer literature had extrapolated, on the basis of accelerated ageing tests in the laboratory (Cole, 2000), that a slight rise in temperature could double failure rates, Pinheiro et al. (2007) demonstrated a far more nuanced reality in the field. Within moderate operating ranges, high temperatures show no significant correlation with increased failure rates, except to a very limited extent on very old hard drives (3 to 4 years in use). Surprisingly, it is very low temperatures (below 30 °C) that are found to be correlated with higher AFR (Annualised Failure Rate) rates (Pinheiro, Weber & Barroso, 2007). Similarly, the overall utilisation rate (workload/duty cycle) exhibits a complex effect: heavy workloads affect the drive only very moderately after its first year of service, which is explained by a ‘survival of the fittest’ phenomenon whereby units prone to teething problems are rapidly weeded out during the initial phase of infant mortality.

2.1.2 Critical S.M.A.R.T. Health Metrics

The Google study scientifically validated the existence of four S.M.A.R.T. parameters that are highly predictive of imminent failures, for which the first occurrence (critical proactivity threshold of 1) leads to a dramatic increase in the probability of short-term failure (within 60 days):

- **Scan Errors:** Physical errors detected during background surface scans. After the first occurrence, the probability of a drive failing increases 39-fold within 60 days (Pinheiro, Weber & Barroso, 2007).
- **Reallocation Counts:** Represents the number of defective sectors detected and dynamically reallocated to healthy physical reserve areas. The occurrence of the first reallocation increases the risk of failure by a factor of 14 within 60 days (Pinheiro, Weber & Barroso, 2007).
- **Offline Reallocations:** Reallocations of defective sectors identified solely during background cleaning routines (offline scrubbing). The risk of failure is then multiplied by 21 (Pinheiro, Weber & Barroso, 2007).

- **Probational Counts:** Indicates the number of suspect sectors temporarily placed on probation. This intermediate signal is also critical and increases the probability of imminent failure by a factor of 16 (Pinheiro, Weber & Barroso, 2007).

Nevertheless, the key lesson from this study for any machine learning project lies in highlighting the intrinsic limitations of raw SMART telemetry. Indeed, 56 per cent of hard disk drives that suffered an actual failure in production had absolutely no anomalies recorded for any of these four critical indicators prior to their failure. Furthermore, 36 per cent of all failing drives showed no alerts across any of the monitored SMART parameters (Pinheiro, Weber & Barroso, 2007). SMART cannot therefore, on its own, constitute a simple deterministic prediction solution, highlighting the need to employ approaches for detecting non-linear weak signals and drift time windows.

To organise these observations, Table 1 summarises the quantitative impact of the SMART metrics according to Pinheiro et al. (2007):

S.M.A.R.T. Attribute	Metric Name	Primary Physical Meaning	Failure Risk Multiplier (60 Days)
S.M.A.R.T. 197	Current Pending Sector Count	Count of unstable sectors awaiting reallocation	39x
S.M.A.R.T. 5	Reallocated Sectors Count	Count of damaged physical sectors remapped	14x
S.M.A.R.T. 198	Offline Uncorrectable Sector	Uncorrectable errors during background scan	21x
S.M.A.R.T. 196	Reallocation Event Count	Total number of remapping operations	16x

Table 1: *Quantitative impact of critical S.M.A.R.T. metrics on 60-day failure probability (adapted from Pinheiro et al., 2007).*

2.2 The challenge of silent corruption and disk typologies

Beyond the outright physical failures studied by Google, silent data corruption poses a major threat to operational reliability. The NetApp study (Bairavasundaram et al., 2008), based on more than 400,000 cases of checksum mismatches analysed over a period of 41 months, revealed profoundly distinct behavioural patterns depending on the class of equipment deployed. Storage infrastructures generally coexist with two families of physical drives: Nearline (SATA) drives and Enterprise (FC – Fibre Channel) drives.

2.2.1 Drive types and asymmetric behaviour in response to anomalies

NetApp’s empirical study demonstrates that Nearline (SATA) drives and their adaptors develop silent corruption at a rate an order of magnitude higher than Enterprise-class (FC) drives. During the first 17 months of operation, the probability of at least one checksum mismatch occurring stands at 0.66 per cent for nearline drives, compared with just 0.06 per cent for enterprise-class drives (Bairavasundaram et al., 2008). Nevertheless, a paradox emerges when analysing the behaviour of the affected population: once an Enterprise-class drive develops its first anomaly, it generates, on average, a cumulative volume of errors far greater than that of a faulty Nearline drive (a median of 10 corruption events per affected Enterprise drive, compared with 2 for SATA). This highly unstable behaviour of high-end components justifies, according to the authors, a strict on-call policy: the immediate and systematic physical replacement of an Enterprise drive as soon as the first anomaly is detected, whereas Nearline drives can be managed on a scheduled or batch basis (Bairavasundaram et al., 2008).

2.2.2 Physical localisation and risks of dual failure

Checksum mismatches exhibit extremely pronounced spatial and temporal localisation. On the one hand, the vast majority of corruption incidents on an affected drive are concentrated in contiguous sectors (average run length of 3.4 consecutive blocks). This very tightly clustered spatial locality indicates that corruption generally occurs during batch write operations or torn writes (Bairavasundaram et al., 2008). This behaviour differs from that of latent sector errors (LSEs), which are caused by physical scratches or media defects that extend over a much wider spatial radius (of the order of 10 MB) and span several physical tracks. Over time, autocorrelation remains very strong over periods of several months (up to 10 months), indicating that anomalies tend to accumulate in successive bursts. The study highlights a critical risk related to system engineering: approximately 8 per cent of nearline disk failures are only detected during the critical RAID reconstruction phase, following the failure of the first disk in the array (Bairavasundaram et al., 2008). During this phase, the absence of proactive prediction and aggressive scrubbing exposes systems to double disk failures, leading to catastrophic data loss. This phenomenon highlights the vital importance of implementing a predictive model capable of intervening prior to this high-risk operational disruption phase.

2.3 Mathematical and methodological theoretical framework in data science

The design of a predictive model with a time horizon of 14 to 30 days, in line with the research question of this study, requires the rigorous application and justification of three theoretical pillars of data science: the management of extreme class imbalance, the choice of the mathematical formulation of the problem, and theories of algorithmic explainability.

2.3.1 Handling Extreme Class Imbalance (Imbalanced Data)

In a data centre production environment, hard disk failures are a rare occurrence, affecting less than 1–2 per cent of disks per year (Bairavasundaram et al., 2008; Pinheiro, Weber & Barroso, 2007). This massive class imbalance poses a major mathematical challenge for machine learning. A naive classifier that systematically predicts the majority class (‘Healthy Disk’) would achieve an overall accuracy of over 98 per cent, whilst being completely ineffective at detecting failures. To overcome this issue, the data science literature recommends the use of advanced sampling techniques. These include, in particular, oversampling of the minority class via the synthesis of synthetic neighbours (SMOTE – Synthetic Minority Over-sampling Technique), controlled downsampling of the majority class (random downsampling of the ‘Healthy’ group) or the incorporation of class weights into robust algorithms such as XGBoost or Random Forest.

2.3.2 Problem formulation: Time-window classification vs. survival analysis

Two complementary methodological approaches compete when formulating predictions over time: * Time-Window Binary Classification: This method involves discretising the time domain by defining a binary label $Y = 1$ if the disc suffers an actual failure within the (N)-day projection window (in our case, 14 to 30 days), and $Y = 0$ otherwise. This allows the input to standard classification models to be standardised, but results in a loss of information regarding the temporal drift dynamics of the telemetry data.

- Survival Analysis (Cox Model / Time-to-Failure): Borrowed from reliability engineering, this mathematical approach seeks to directly predict the ‘time to failure’ (TTF). It offers the major theoretical advantage of being able to rigorously handle right-censored data, i.e. disks removed from production for other reasons (upgrades or maintenance) before they have actually suffered a physical failure.

2.3.3 Assessing the importance of variables and explainability (XAI)

In order to identify the most critical physical variables (addressing the second part of our research question), it is necessary to look beyond algorithmic black boxes. Whilst intrinsic importance measures such

as Gini impurity (used by random forests) or information gain provide an initial overall classification, they prove limited when faced with complex and non-linear interactions. This is why modern theory on software explainability employs approaches based on cooperative game theory. SHAP (SHapley Additive exPlanations) values and LIME models make it possible to mathematically isolate the exact marginal impact of each SMART attribute on the prediction of the probability of failure at each time ‘t’. This transparency is essential for justifying and legitimising the proactive replacement of a hard drive to SRE teams.

2.4 Systemic convergence: from raw hardware telemetry to sre observability

The ultimate value of this research project lies in the integration of this predictive model within the modern conceptual framework of observability engineering and SRE practices. The transition from passive hardware monitoring to proactive, AIOps-driven observability is redefining the management of software infrastructure reliability.

2.4.1 The epistemological shift: from traditional monitoring to proactive observability

Traditional monitoring works by checking predefined rules and applying static thresholds (for example, triggering an alert if one of the critical SMART parameters rises above zero). Majors, Fong-Jones and Miranda (2022) demonstrate that this paradigm is fundamentally reactive, generates incessant false positives and alert fatigue, and proves incapable of diagnosing complex, previously unencountered behaviours (‘unknown-unknowns’). Conversely, observability is defined as the ability to infer the internal state of a complex system solely by analysing its external outputs based on structured events with high cardinality and high dimensionality. From this perspective, SMART telemetry data should no longer be regarded as isolated, low-level alerts, but as continuous infrastructure metrics. The machine learning model thus acts as an AIOps (AI for IT Operations) layer: by correlating temporal trends and variations in weak, non-linear signals, it extracts predictive intelligence where a human or a conventional threshold would see only background noise.

2.4.2 Operational impact: eliminating Toil and protecting SLOs

The operational implementation of this prediction aligns rigorously with the SRE pillars theorised by Google:

- **Elimination of ‘Toil’:** ‘Toil’ refers to all manual, repetitive, automatable operational tasks that offer no lasting added value, the volume of which grows linearly with the size of the system (Beyer et al., 2016). Predicting a hard disk failure 14 to 30 days in advance breaks this linearity: hardware replacement is planned and systematically integrated into the physical teams’ backlog during normal maintenance windows, eliminating emergency interventions and reactive night-time on-call duties (Beyer et al., 2018).
- **Protection of SLOs and Error Budgets:** When a physical drive suffers a sudden failure in production, the distributed storage infrastructure (Hadoop, Ceph, or Colossus at Google) must immediately initiate data reconstruction using geographical replicas or erasure coding. This reconstruction process consumes a considerable amount of network bandwidth and saturates the I/O performance of adjacent disks. For the end user, this overhead results in a drastic increase in application query latency, which degrades the user experience and directly depletes the latency and availability error budgets set by the SLOs. Thanks to proactive prediction, the orchestrator or SRE team can switch the unit to a phased-out mode (Lame Duck mode or ‘drain data’ mode). Data is thus migrated seamlessly, in a planned manner and completely transparently to the user, protecting the organisation’s error budget (Thurgood et al., 2018).
- **Addressing alert fatigue:** SRE is constantly striving to eliminate non-actionable alerts. A well-calibrated machine learning model, which prioritises a rigorous balance between precision (to avoid wasting healthy hardware and false alarms) and recall (to ensure no actual failures are missed), guarantees that every alert issued corresponds to a legitimate and actionable hardware anomaly (Beyer et

al., 2018). In conclusion, the integration of a SMART predictive hardware telemetry tool lies at the heart of current challenges in reliability engineering research. By transforming weak physical signals into a proactive, actionable prediction spanning 14 to 30 days, this work demonstrates how data science and AIOps come together to preserve the large-scale resilience of modern distributed systems.

3 Methods and Analysis

This chapter presents the end-to-end Machine Learning pipeline developed for predictive maintenance, detailing the data ingestion workflow, feature engineering transformations, down-sampling strategies, and model architecture. The framework is designed for high-dimensional and highly imbalanced telemetry, with the objective of forecasting drive failures within a proactive maintenance window while minimizing false alarms and unnecessary operational interventions. Following established literature (Pinheiro et al.; Schroeder & Gibson), raw drive telemetry was filtered to focus on 5 critical hardware failure predictors: SMART 5 (Re-allocated Sectors), 187 (Reported Uncorrectable Errors), 196 (Reallocation Events), 197 (Current Pending Sectors), and 198 (Offline Uncorrectable Sectors).

3.1 Methodological Framework and Experimental Design

To address the challenges of predicting hardware degradation in production environments, our analytical approach rests on three core methodological pillars:

1. **Temporal Binary Classification Framing ($H = 14$):** Predictive maintenance is formulated as a daily binary classification task. For any given drive i on day t , the target variable $Y_{it} \in \{0, 1\}$ indicates whether the drive will suffer a catastrophic failure within the look-ahead horizon $H = 14$ days ($[t + 1, t + 14]$). This 14-day lead time reflects Site Reliability Engineering (SRE) operational requirements, providing sufficient time for automated data migration without triggering emergency maintenance.
2. **Prevention of Temporal Data Leakage:** Because S.M.A.R.T. metrics represent time-series sequences generated by physical media, standard cross-validation techniques (such as random k-fold splitting) would introduce severe temporal data leakage between contiguous daily snapshots of the same drive. To ensure strict evaluation integrity, all datasets are split temporally: historical data forms the training partition, while subsequent temporal blocks are reserved exclusively for validation and testing.
3. **“Design-by-Constraint” Engineering (6 GB RAM Limit):** Operating under a strict 6 GB Virtual Machine RAM footprint requires a memory-aware execution pipeline. Rather than loading multi-gigabyte quarterly datasets into memory in their entirety, our methodology uses stream-processing via `data.table::fread`, selective column indexing, down-casting of data types, and proactive garbage collection (`gc()`) at key pipeline boundaries.

3.2 Data Ingestion, Cleaning, Label Alignment and Memory Management Strategy

The ingestion pipeline processes Backblaze telemetry through selective column parsing and vectorized memory management to ensure reproducible execution within the 6 GB RAM environment. The dataset consists of daily S.M.A.R.T. (Self-Monitoring, Analysis, and Reporting Technology) attribute logs collected across large enterprise drive populations. Data cleaning involves filtering out drive models with insufficient sample sizes, imputing missing sensor readings via temporal forward filling, and removing invariant operational metrics.

3.2.1 Ingestion and Standardization of Backblaze Telemetry Data

The ingestion module implements an automated workflow for downloading, extracting, and standardizing Backblaze’s quarterly telemetry archives. Its purpose is to guarantee reproducible data acquisition while maintaining strict memory efficiency within the 6 GB RAM execution environment. The pipeline relies on

selective parsing and vectorized operations to avoid unnecessary memory overhead, a critical requirement when processing high volume daily S.M.A.R.T. logs across large enterprise drive fleets.

The script begins by configuring a dedicated raw directory and retrieving the official Backblaze Q1-2024 dataset. To ensure robustness, the download step includes a file integrity check and forces the use of the libcurl engine, which bypasses R's default 60 second timeout limitations. Once the archive is available locally, the pipeline performs controlled decompression and initiates a structural normalization phase.

A recurring issue with Backblaze archives is the presence of macOS specific artefacts—such as the `__MACOSX` folder and hidden Apple cache files—that can disrupt downstream parsing. The ingestion script systematically removes these artefacts, flattens nested directory structures, and consolidates all valid CSV telemetry files into a unified folder. This ensures that subsequent processing stages—feature engineering, temporal alignment, and failure label construction—operate on a clean and predictable file layout.

The final validation step enumerates all daily telemetry files matching the expected naming convention (e.g., `2024_01_01.csv`). This provides a deterministic set of input files for the analytical pipeline and eliminates the risk of silent ingestion errors. By enforcing structural consistency at this early stage, the pipeline ensures that downstream transformations—such as feature engineering, temporal imputation, and failure window labeling—can be applied reliably across the entire dataset.

```
local({
  library(httr)
  options(timeout = 3600) #Global R timeout setting (1 hour = 3,600 seconds)

  # 1. Directory configuration
  raw_dir <- "data/raw"
  if (!dir.exists(raw_dir)) dir.create(raw_dir, recursive = TRUE)

  # 2. Backblaze's official Q1 2024 URL
  zip_url <- "https://f001.backblazeb2.com/file/Backblaze-Hard-Drive-Data/data_Q1_2024.zip"
  zip_path <- file.path(raw_dir, "data_Q1_2024.zip")

  # 3. Download if the archive does not exist or is incomplete
  if (!file.exists(zip_path) || file.info(zip_path)$size < 100000000) {
    message("--> Download the Backblaze Q1 2024 archive (please wait)...")
    #Force the use of the libcurl/curl engine to bypass the RStudio driver, which sets the
    ↪ timeout to 60 seconds
    download.file(
      url      = zip_url,
      destfile = zip_path,
      method   = if (capabilities("libcurl")) "libcurl" else "auto",
      mode     = "wb"
    )
  } else {
    message("--> Zip archive already present locally.")
  }

  # 4. Extracting the archive
  message("--> Extract the zip archive to data/raw/...")
  unzip(zip_path, exdir = raw_dir)

  # 5. Resolving the issue of unwanted macOS subfolders and files
  message("--> Standardisation of the structure of CSV files...")

  # A. Immediate deletion of the __MACOSX junk folder created by macOS
  macosx_dir <- file.path(raw_dir, "__MACOSX")
})
```

```

if (dir.exists(macosx_dir)) {
  unlink(macosx_dir, recursive = TRUE)
  message("    [OK] The __MACOSX folder has been deleted.")
}

# B. Migrate the CSV files from the extracted subfolder to data/raw/
nested_dir <- file.path(raw_dir, "data_Q1_2024")
if (dir.exists(nested_dir)) {
  nested_csvs <- list.files(nested_dir, pattern = "\\*.csv$", full.names = TRUE)
  if (length(nested_csvs) > 0) {
    file.rename(nested_csvs, file.path(raw_dir, basename(nested_csvs)))
    message(sprintf("    [OK] %d CSV files moved from the subfolder to %s/",
      ↪ length(nested_csvs), raw_dir))
  }
  # Deleting the empty folder
  unlink(nested_dir, recursive = TRUE)
}

# C. Cleaning up Apple's hidden files (ex: ._2024-01-01.csv)
apple_dot_files <- list.files(raw_dir, pattern = "\\._.*\\.csv$", full.names = TRUE)
if (length(apple_dot_files) > 0) {
  file.remove(apple_dot_files)
  message(sprintf("    [OK] %d malicious Apple cache files (._*.csv) deleted.",
    ↪ length(apple_dot_files)))
}

# 6. Final check
valid_csvs <- list.files(raw_dir, pattern = "^2024-.*\\.csv$", full.names = TRUE)
message(sprintf("==> Import completed successfully: %d valid CSV files ready for
  ↪ processing.", length(valid_csvs)))
})
gc(verbose = FALSE)

```

```

##          used  (Mb) gc trigger  (Mb) max used  (Mb)
## Ncells 2632329 140.6   4705135 251.3  4705135 251.3
## Vcells 4533925  34.6   10146329  77.5  6570417  50.2

```

Memory Footprint and Computational Efficiency Note:

The garbage collection summary reported above provides a fine-grained view of the pipeline's memory footprint by detailing R's dual-allocation mechanism:

- **Ncells (Node cells):** Represent the structural backbone and administrative bookkeeping of the R environment. They store metadata, symbols, and pointers—conceptually acting like the library catalog that indexes where objects are located and how they are linked, rather than holding the raw content itself.
- **Vcells (Vector cells):** Represent the actual memory allocated to atomic vectors, which constitute the core raw data payload (such as numerical arrays, logical flags, and character strings). In simple terms, if Ncells serve as the architectural blueprint and references, Vcells store the physical values themselves.

Here, Vcells used accounts for 34.6 MB (with a historical peak of 53.3 MB), while Ncells used occupies 141.5 MB (peak at 251.0 MB), resulting in a modest cumulative footprint of approximately 176.1 MB. By leveraging stream-processing via `data.table::fread` and executing targeted `gc()` calls post-ingestion, the pipeline prevents memory fragmentation across both structural pointers and primary data vectors, ensuring predictable performance well within the enforced 6 GB Virtual Machine RAM limit.

3.2.2 From Raw Telemetry to Analytical Readiness

Before proceeding to exploratory analysis, the ingestion pipeline ensures that all telemetry is structurally consistent, memory-efficient, and temporally aligned. The removal of corrupted artefacts, consolidation of daily CSV snapshots, and normalization of S.M.A.R.T. fields provide a clean analytical foundation.

With these ingestion and normalization procedures in place, the dataset is memory-optimized and ready for empirical inspection. The next phase of our methodology carries out an Exploratory Data Analysis (EDA) to assess the underlying distribution of S.M.A.R.T. metrics, diagnose severe class imbalance, and validate empirical failure signatures prior to feature engineering and model training.

4 Exploratory Data Analysis (EDA) of Backblaze Telemetry

Exploratory Data Analysis provides an initial statistical and operational understanding of Backblaze telemetry. This chapter examines global failure patterns, S.M.A.R.T. feature behaviour, temporal trends, and hardware heterogeneity across drive models and capacities. These diagnostics reveal the structural challenges of the dataset—extreme class imbalance, non-linear degradation signals, and feature redundancy—and guide the design of the predictive pipeline.

4.1 Global Failure Patterns and Feature Behaviour

The exploratory analysis module provides a structured examination of the processed telemetry, enabling a first assessment of class imbalance, feature behaviour, and temporal failure dynamics. The first diagnostic confirms the extreme imbalance between normal operational days and failure-window observations, with failure rates below 0.1%—a typical characteristic of enterprise-grade storage fleets.

A second analysis examines failure rates by drive model, revealing substantial heterogeneity across hardware families. Several high-volume models exhibit disproportionately elevated failure frequencies, motivating model-aware stratification during training.

The module then compares S.M.A.R.T. feature distributions between healthy drives and drives approaching failure. Indicators such as S.M.A.R.T. 5 (Reallocated Sectors), S.M.A.R.T. 187 (Uncorrectable Errors), and S.M.A.R.T. 197 (Pending Sectors) exhibit step-function jumps immediately prior to termination, demonstrating strong discriminative potential. Conversely, cumulative counters such as S.M.A.R.T. 9 (Power-On Hours) act primarily as confounding variables rather than direct failure signals.

A Spearman correlation heatmap further evaluates redundancy and collinearity among raw S.M.A.R.T. metrics, supporting feature-selection decisions and reducing dimensionality. Temporal failure-rate trends and failure rates by storage capacity complete the diagnostic view, contextualising reliability across time and hardware classes.

```
local({
  # 1. Loading the required packages
  required_packages <- c("data.table", "tidyverse", "lubridate", "ggplot2")

  for (pkg in required_packages) {
    if (!require(pkg, character.only = TRUE)) {
      install.packages(pkg)
      library(pkg, character.only = TRUE)
    }
  }

  # 2. Configuring memory and environment options
  options(scipen = 999) # Disable scientific notation
  gc() # Freeing up RAM (Garbage Collection)

  # 3. Defining relative paths
  PATH_RAW <- "data/raw/"
  PATH_PROCESSED <- "data/processed/"
  PATH_FIGURES <- "outputs/figures/"

  # 4. Defining quarter to analyse
  DATASET_TAG <- "Q1_2024" # Adjust if processing a different quarter (e.g., "Q2_2024")
})
```

```

# Dynamic output file construction based on dataset tag
OUTPUT_FILENAME <- sprintf("dt_processed_%s.rds", tolower(DATASET_TAG))
OUTPUT_FILEPATH <- file.path(PATH_PROCESSED, OUTPUT_FILENAME)
EDA_RAW_FILEPATH <- file.path(PATH_PROCESSED, "eda_raw_summary.rds")

cat("=====\n")
cat(" Environment successfully initialised for the VM (6 GB RAM)\n")
cat("=====\n")

# =====
# Phase 2: Optimised data ingestion, EDA & feature engineering
# =====

library(data.table)
library(tidyverse)
library(lubridate)

# 1. Selection of strategic columns (RAM optimisation)
# -----
# Columns identified in the literature review (Pinheiro et al., 2007):
# SMART 5: Reallocated Sectors Count
# SMART 187: Reported Uncorrectable Errors
# SMART 188: Command Timeout
# SMART 197: Current Pending Sector Count
# SMART 198: Offline Uncorrectable / Scan Errors
# SMART 9: Power-On Hours (Drive Age)
# SMART 194: Temperature (For monitoring/validation)

keep_cols <- c(
  "date", "serial_number", "model", "capacity_bytes", "failure",
  "smart_5_raw", "smart_9_raw", "smart_187_raw", "smart_188_raw",
  "smart_194_raw", "smart_197_raw", "smart_198_raw"
)

# 2. Optimised data ingestion with data.table
# -----
cat("Starting to import the CSV files...\n")

# List all the CSV files extracted to the data/raw/ directory
csv_files <- list.files(path = PATH_RAW, pattern = "*.csv", full.names = TRUE)

if (length(csv_files) == 0) {
  stop("Please note: No CSV files were found in the 'data/raw/' directory. Please place
  ↳ the Backblaze data there.")
}

# Fast ingestion and column selection only
dt_list <- lapply(csv_files, function(file) {
  fread(
    file,
    select = keep_cols,
    colClasses = c(capacity_bytes = "numeric"), # Force dual-channel/64-bit reading
  )
})

```



```

    showProgress = FALSE
  )
})

# Fast ingestion and column selection only
dt_raw <- rbindlist(dt_list, fill = TRUE)
rm(dt_list) # Immediate memory release
gc()

cat(sprintf("Input complete: %s rows and %s columns loaded into memory.\n",
            format(nrow(dt_raw), big.mark = " "), ncol(dt_raw)))

# Extraction & Saving Raw EDA Metadata
raw_failure_dist <- dt_raw[, .N, by = failure][order(failure)]
top_10_models <- dt_raw[, .N, by = model][order(-N)][1:10]

# Interception and explicit saving of raw metrics
eda_summary <- list(
  failure_dist = raw_failure_dist,
  top_models = top_10_models,
  total_rows = nrow(dt_raw)
)
saveRDS(eda_summary, file = EDA_RAW_FILEPATH)

# 3. Initial quick review and quality control (basic EDA)
# -----
cat("\n--- Distribution of the target variable (Failure) ---\n")
print(table(dt_raw$failure, useNA = "ifany"))

cat("\n--- Top 10 most common disc models ---\n")
top_models <- dt_raw[, .N, by = model][order(-N)][1:10]
print(top_models)

# Memory clean-up
gc

# =====
# 4. Data Filtering & Model Selection
# =====
cat("\n--- Filtering top drive models to reduce noise ---\n")

# Keep top drive models representing majority of the population
top_model_names <- top_10_models[1:5, model]
dt_filtered <- dt_raw[model %in% top_model_names]

# Free up memory from raw data
rm(dt_raw)
gc()

# Sort data chronologically per drive (essential for time-series feature engineering)
setkey(dt_filtered, serial_number, date)

# =====

```

```

# 5. Temporal Feature Engineering (SRE Prediction Horizon Y_it)
# =====
cat("\n--- Creating 14-day Failure Horizon (Target Y_it) ---\n")

# Define target window: 14-day lead time for proactive replacement
PREDICTION_WINDOW_DAYS <- 14

# Calculate max failure date per drive serial number
dt_filtered[, max_failure := max(failure), by = serial_number]
dt_filtered[, failure_date := as.IDate(ifelse(failure == 1, date, NA)), by =
  ↪ serial_number]
dt_filtered[, failure_date := min(failure_date, na.rm = TRUE), by = serial_number]

# Define binary label Y_it: 1 if drive fails within [t, t + 14 days]
dt_filtered[, days_to_failure := as.numeric(failure_date - date)]
dt_filtered[, target_14d := ifelse(!is.na(days_to_failure) &
  ↪ days_to_failure >= 0 &
  ↪ days_to_failure <= PREDICTION_WINDOW_DAYS, 1, 0)]

# Clean temporary calculation columns
dt_filtered[, c("max_failure", "failure_date", "days_to_failure") := NULL]

# =====
# 6. Feature Engineering: SMART Deltas (7-day Rate of Change)
# =====
cat("\n--- Computing 7-day SMART Delta Features ---\n")

# Calculate 7-day changes for critical SMART metrics (Pinheiro et al., 2007)
# Delta = Current raw value - Value 7 days ago
dt_filtered[, smart_5_delta7 := smart_5_raw - shift(smart_5_raw, 7, type = "lag"), by =
  ↪ serial_number]
dt_filtered[, smart_187_delta7 := smart_187_raw - shift(smart_187_raw, 7, type = "lag"),
  ↪ by = serial_number]
dt_filtered[, smart_197_delta7 := smart_197_raw - shift(smart_197_raw, 7, type = "lag"),
  ↪ by = serial_number]

# Impute initial NA values resulting from lag calculations with 0
na_cols <- c("smart_5_delta7", "smart_187_delta7", "smart_197_delta7")
for (j in na_cols) set(dt_filtered, which(is.na(dt_filtered[[j]])), j, 0)

# =====
# 7. Summary & Export Processed Dataset
# =====
cat("\n--- Target Horizon Distribution (14-day window) ---\n")
print(table(dt_filtered$target_14d, useNA = "ifany"))

# Save processed table to disk
cat(sprintf("\nSaving processed dataset to %s...\n", OUTPUT_FILEPATH))

# Using 'xz' compression for optimal RAM & Disk efficiency on VM
saveRDS(dt_filtered, file = OUTPUT_FILEPATH, compress = "xz")

```

```
cat("Processing completed successfully!\n")
gc()
})
gc(verbose = FALSE)
```

Global Failure Patterns and Fleet Composition Analysis:

The exploratory analysis reveals three core structural characteristics of the Backblaze telemetry dataset across its raw and processed stages:

- **Extreme Class Imbalance:** Hard drive failures (`failure = 1`) represent a tiny fraction of total operational records—exactly 978 failures out of 25,189,213 daily logs ($\approx 0.0039\%$), as detailed in Table 1. When extended to the 14-day prediction horizon (Y_{it}), positive failure instances increase to 5,247 records (0.0398%) within the filtered Top 5 fleet (Table 4). This severe class imbalance dictates specialized resampling or cost-sensitive learning algorithms during downstream model formulation to avoid classifier bias toward the majority class.
- **Fleet Homogeneity and Model Concentration:** As detailed in Table 2, the hard drive infrastructure is heavily concentrated around a few primary drive families. Toshiba models (MG07ACA14TA and MG08ACA16TA) together constitute over 26.77% of the overall telemetry population, followed by high-capacity Seagate (ST16000NM001G), Western Digital (WDC WUH721816ALE6L4), and Seagate (ST12000NM0008) units. Together, these Top 5 models represent 52.29% of the global dataset (13,172,916 daily logs). Filtering out fringe, low-count drive models beyond this Top 5 is thus an essential step to mitigate noise and prevent model overfitting to rare hardware architectures. Within this filtered Top 5 fleet (Table 3), Toshiba models (MG07ACA14TA and MG08ACA16TA) alone constitute 51.19% of the operational subset (26.25% and 24.94% respectively).
- **SMART Feature Heterogeneity and Multicollinearity:** As illustrated by the distribution profiles in Figure 3 (in Section 4.1), raw SMART attributes exhibit highly non-linear, zero-inflated behaviors with extreme scale disparities. Critical failure indicators - such as `smart_5` (Reallocated Sectors), `smart_187` (Uncorrectable Errors), and `smart_197` (Pending Sectors) - remain completely flat at zero for healthy drives but display severe continuous spikes for drives failing within the 14-day horizon, whereas operational baselines like `smart_194` (Temperature) follow symmetric distributions across both classes. Furthermore, the correlation structure in Figure 4 (in Section 4.1) isolates the primary degradation cluster, notably showing moderate collinearity between sector metrics (`smart_197_raw` and `smart_198_raw` at 0.49) alongside low linear redundancy with baseline controls (`smart_9_raw` (Power-On Hours) and `smart_194_raw`). This validates both the $\log(1+x)$ transformations/scaling applied during preprocessing and the retention of temporal rate-of-change features (`delta7`).

The quantitative distributions and fleet statistics supporting these observations are presented in Tables 1 through 4 on the following page.

Table 1: Raw Target Variable Distribution (25.1M Daily Logs)

Status (failure)	Description	Observations (N)	Proportion (%)
0	Operational (Healthy)	25 188 235	99.9961 %
1	Failed	978	0.0039 %

Table 2: Top 10 Most Common Hard Drive Models (Raw Telemetry)

Drive Model	Observations (N)	Share of Total Population (%)
TOSHIBA MG07ACA14TA	3 458 359	13.73 %
TOSHIBA MG08ACA16TA	3 285 527	13.04 %
ST16000NM001G	2 579 015	10.24 %
WDC WUH721816ALE6L4	2 083 984	8.27 %
ST12000NM0008	1 766 031	7.01 %
ST8000NM0055	1 262 294	5.01 %
HGST HUH721212ALE604	1 198 593	4.76 %
ST12000NM001G	1 195 601	4.75 %
ST14000NM001G	974 767	3.87 %
HGST HUH721212ALN604	958 729	3.81 %

Table 3: Selected Top 5 Fleet Composition (Processed Dataset)

Drive Model (Top 5)	Observations (N)	Share of Filtered Fleet (%)
TOSHIBA MG07ACA14TA	3 458 359	26.25 %
TOSHIBA MG08ACA16TA	3 285 527	24.94 %
ST16000NM001G	2 579 015	19.58 %
WDC WUH721816ALE6L4	2 083 984	15.82 %
ST12000NM0008	1 766 031	13.41 %

Table 4: Transformed Target Horizon Distribution (Filtered Fleet)

Target (Y_it)	Description	Observations (N)	Proportion (%)
0	No Failure in Horizon	13 167 669	99.9602 %
1	Failure Predicted in Window	5 247	0.0398 %

4.2 Extended Exploratory Analysis: Correlation, Temporal Dynamics, and Capacity Effects

A complementary analysis deepens the statistical examination of failure dynamics by quantifying patterns across models, capacities, temporal windows, and key S.M.A.R.T. indicators. This module reinforces the initial findings and provides additional operational insights.

The correlation analysis identifies clusters of redundant S.M.A.R.T. indicators, informing regularisation and feature-selection strategies. Temporal failure-rate trends reveal operational fluctuations and potential seasonal effects, highlighting the importance of temporal alignment in predictive maintenance pipelines. Failure rates by storage capacity show how reliability varies across drive sizes, validating model generalisation across heterogeneous hardware populations.

Together, these analyses form a comprehensive diagnostic foundation for the predictive pipeline. They validate statistical assumptions, guide feature-engineering decisions, and ensure that subsequent modelling steps are grounded in empirical evidence extracted from the dataset.

```
local({
  # =====
  # Script: 01b_exploratory_data_analysis.R
  # Project: Predictive Hard Drive Failure Modelling for SRE Operations
  # Description: Exploratory Data Analysis (EDA) & Feature Distributions
  # =====

  # 0. Load Libraries
  required_packages <- c("data.table", "tidyverse", "lubridate", "scales", "gridExtra")

  for (pkg in required_packages) {
    if (!require(pkg, character.only = TRUE)) {
      install.packages(pkg)
      library(pkg, character.only = TRUE)
    }
  }

  options(scipen = 999)
  theme_set(theme_minimal(base_size = 12))

  # 1. Directories Setup
  PATH_PROCESSED <- "data/processed/"
  PATH_FIGS <- "outputs/figures/"

  if (!dir.exists(PATH_FIGS)) dir.create(PATH_FIGS, recursive = TRUE)

  INPUT_FILEPATH <- file.path(PATH_PROCESSED, "dt_processed_q1_2024.rds")

  cat("Loading processed dataset for EDA...\n")
  dt <- readRDS(INPUT_FILEPATH)

  # =====
  # 2. EDA 1: Target Class Imbalance
  # =====
  cat("\n--- Generating Figure 1: Target Distribution ---\n")

  target_counts <- dt[, .(Count = .N), by = target_14d]
  target_counts[, Percent := (Count / sum(Count)) * 100]
```

```

target_counts[, Label := ifelse(target_14d == 1, "Failure Window (14d)", "Healthy /
  ↳ Normal")]

p1 <- ggplot(target_counts, aes(x = Label, y = Count, fill = Label)) +
  geom_bar(stat = "identity", width = 0.5, show.legend = FALSE) +
  geom_text(aes(label = sprintf("%s\n(%.3f%%)", format(Count, big.mark = " "), Percent)),
    ↳
      vjust = -0.3, size = 4, fontface = "bold") +
  scale_y_continuous(labels = comma, expand = expansion(mult = c(0, 0.15))) +
  scale_fill_manual(values = c("#2c3e50", "#e74c3c")) +
  labs(
    title = "Class Imbalance in Hard Drive Failure Dataset (Q1 2024)",
    subtitle = "Extreme imbalance between operational days and 14-day failure windows",
    x = "Class Label",
    y = "Number of Daily Observations"
  )

ggsave(file.path(PATH_FIGS, "fig1_class_imbalance.png"), plot = p1, width = 8, height =
  ↳ 5, dpi = 300)

# =====
# 3. EDA 2: Failure Rate by Drive Model
# =====
cat("--- Generating Figure 2: Failure Rate by Drive Model ---\n")

# cleanup
dt_clean_model <- dt[!is.na(model) & trimws(as.character(model)) != ""]

# statistics calculation
model_stats <- dt_clean_model[, .(
  Total_Obs = .N,
  Failures = sum(as.numeric(failure), na.rm = TRUE),
  Failure_Rate_pct = (sum(as.numeric(failure), na.rm = TRUE) / .N) * 100
), by = model][order(-Total_Obs)]

# Keep at most the top 10 models for visualization clarity
top_n_models <- min(10, nrow(model_stats))
model_stats_top <- model_stats[1:top_n_models]
model_stats_top[, model := factor(model, levels = model[order(Failure_Rate_pct)])]

# Dynamic subtitle based on the number of detected models
sub_title_f2 <- sprintf("Daily failure rates for top %d models by volume (out of %d total
  ↳ models)",
  top_n_models, nrow(model_stats))

p2 <- ggplot(model_stats_top, aes(x = model, y = Failure_Rate_pct)) +
  geom_col(fill = "#3498db", width = 0.6) +
  geom_text(aes(label = sprintf("%.4f%%", Failure_Rate_pct)), hjust = -0.1, size = 3.5,
    ↳ fontface = "bold") +
  coord_flip() +
  scale_y_continuous(expand = expansion(mult = c(0, 0.25)), labels = function(x)
    ↳ paste0(sprintf("%.4f", x), "%")) +
  labs(
    title = "Hard Drive Failure Rate by Model",

```

```

    subtitle = sub_title_f2,
    x = "Drive Model",
    y = "Daily Failure Rate (%)"
  )

# Dynamic adjustment of the image height when there are many models
fig_height <- max(4.5, top_n_models * 0.45)
ggsave(file.path(PATH_FIGS, "fig2_failure_rate_by_model.png"), plot = p2, width = 9,
  ↪ height = fig_height, dpi = 300)
rm(dt_clean_model, model_stats, model_stats_top)

# =====
# 4. EDA 3: S.M.A.R.T. Feature Distributions (Healthy vs Failing)
# =====
cat("--- Generating Figure 3: S.M.A.R.T. Feature Distributions ---\n")

# Sample for plotting efficiency
set.seed(42)
sample_dt <- dt[sample(.N, 100000)]
sample_dt[, Status := ifelse(target_14d == 1, "Failing Within 14 Days", "Healthy")]

p3a <- ggplot(sample_dt, aes(x = Status, y = log1p(smart_5_raw), fill = Status)) +
  geom_boxplot(alpha = 0.7, outlier.size = 0.5) +
  scale_fill_manual(values = c("#e74c3c", "#2ecc71")) +
  theme(legend.position = "none") +
  labs(title = "SMART 5 (Reallocated Sectors)", y = "log1p(Count)")

p3b <- ggplot(sample_dt, aes(x = Status, y = log1p(smart_187_raw), fill = Status)) +
  geom_boxplot(alpha = 0.7, outlier.size = 0.5) +
  scale_fill_manual(values = c("#e74c3c", "#2ecc71")) +
  theme(legend.position = "none") +
  labs(title = "SMART 187 (Uncorrectable Errors)", y = "log1p(Count)")

p3c <- ggplot(sample_dt, aes(x = Status, y = log1p(smart_197_raw), fill = Status)) +
  geom_boxplot(alpha = 0.7, outlier.size = 0.5) +
  scale_fill_manual(values = c("#e74c3c", "#2ecc71")) +
  theme(legend.position = "none") +
  labs(title = "SMART 197 (Pending Sectors)", y = "log1p(Count)")

p3d <- ggplot(sample_dt, aes(x = Status, y = smart_194_raw, fill = Status)) +
  geom_boxplot(alpha = 0.7, outlier.size = 0.5) +
  scale_fill_manual(values = c("#e74c3c", "#2ecc71")) +
  theme(legend.position = "none") +
  labs(title = "SMART 194 (Temperature °C)", y = "Raw Value")

p3_combined <- grid.arrange(p3a, p3b, p3c, p3d, ncol = 2)

ggsave(file.path(PATH_FIGS, "fig3_smart_features_distribution.png"), plot = p3_combined,
  ↪ width = 10, height = 8, dpi = 300)

# =====
# 5. EDA 4: S.M.A.R.T. Features Correlation Heatmap (Memory-Safe)
# =====
cat("--- Generating Figure 4: S.M.A.R.T. Correlation Heatmap ---\n")

```

```

# Targeted list: 5 critical failure attributes + 2 checks + key deltas + target
smart_cols <- c(
  "smart_5_raw", "smart_187_raw", "smart_196_raw", "smart_197_raw", "smart_198_raw", #
  ↪ raw critical
  "smart_9_raw", "smart_194_raw", #
  ↪ Controles
  "smart_5_delta7", "smart_187_delta7", "smart_197_delta7" #
  ↪ Deltas
)

# Retain only the columns actually present in the dt dataset
smart_cols <- intersect(smart_cols, names(dt))
#smart_cols <- c("smart_5_raw", "smart_9_raw", "smart_187_raw", "smart_188_raw",
  ↪ "smart_194_raw", "smart_197_raw", "smart_198_raw")

# A sample of 50,000 lines
set.seed(42)
sample_size <- min(50000, nrow(dt))
cor_dt_sample <- dt[sample(.N, sample_size), smart_cols, with = FALSE]

# Cleaning column by column (avoids flattening the data.table)
for (col in smart_cols) {
  if (col != "failure") {
    val <- cor_dt_sample[[col]]
    val[is.na(val) | val < 0] <- 0
    set(cor_dt_sample, j = col, value = log1p(val))
  }
}

# Calculation of the Spearman matrix
cor_mat <- cor(as.matrix(cor_dt_sample), method = "spearman", use =
  ↪ "pairwise.complete.obs")

# Custom 2D-to-long format conversion for ggplot2
cor_df <- expand.grid(Var1 = smart_cols, Var2 = smart_cols, KEEP.OUTATTRS = FALSE)
cor_df$value <- as.vector(cor_mat)
cor_df <- cor_df[!is.na(cor_df$value), ] # Filter out NA values

# Force the order of the axes to match the predefined list
cor_df$Var1 <- factor(cor_df$Var1, levels = smart_cols)
cor_df$Var2 <- factor(cor_df$Var2, levels = rev(smart_cols)) # 'rev' to align the
  ↪ diagonal correctly

p4 <- ggplot(cor_df, aes(x = Var1, y = Var2, fill = value)) +
  geom_tile(color = "white") +
  scale_fill_gradient2(low = "#3498db", mid = "#ffffff", high = "#e74c3c", midpoint = 0,
  ↪ limit = c(-1, 1), name = "Spearman\nCorr") +
  geom_text(aes(label = sprintf("%.2f", value)), color = "black", size = 3) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs(
    title = "S.M.A.R.T. Features Spearman Correlation Matrix",
    subtitle = "Identification of feature redundancy and collinearity among raw
    ↪ indicators (50k sample)",
  )

```



```

    x = "", y = ""
  )

ggsave(file.path(PATH_FIGS, "fig4_correlation_heatmap.png"), plot = p4, width = 8, height
↪   = 6, dpi = 300)
rm(cor_dt_sample, cor_mat, cor_df)

# =====
# 6. EDA 5: Daily Failure Rate Trends Over Time
# =====
cat("--- Generating Figure 5: Temporal Failure Rate Trends ---\n")

daily_trend <- dt[, .(
  Total_Disks = .N,
  Failures = sum(failure),
  Failure_Rate_pct = (sum(failure) / .N) * 100
), by = date][order(date)]

p5 <- ggplot(daily_trend, aes(x = date, y = Failure_Rate_pct)) +
  geom_line(color = "#e74c3c", size = 0.8) +
  geom_smooth(method = "loess", color = "#2c3e50", se = FALSE, linetype = "dashed", span
↪   = 0.3) +
  scale_x_date(date_breaks = "2 weeks", date_labels = "%b %d") +
  scale_y_continuous(labels = function(x) paste0(sprintf("%.4f", x), "%")) +
  labs(
    title = "Daily Hard Drive Failure Rate (Q1 2024)",
    subtitle = "Daily percentage of failing drives with LOESS smoothing trend line",
    x = "Date",
    y = "Daily Failure Rate (%)"
  )

ggsave(file.path(PATH_FIGS, "fig5_temporal_failure_trend.png"), plot = p5, width = 9,
↪   height = 5, dpi = 300)
rm(daily_trend)

# =====
# 7. EDA 6: Failure Rate by Storage Capacity
# =====
cat("--- Generating Figure 6: Failure Rate by Drive Capacity ---\n")

# Filtering out invalid/negative capacities (-1)
dt_cap <- copy(dt[!is.na(capacity_bytes) & capacity_bytes > 0])

# Dynamic conversion to TB
dt_cap[, capacity_tb := round(as.numeric(capacity_bytes) / 1e12)]

# Aggregation by capacity
cap_stats <- dt_cap[capacity_tb > 0, .(
  Total_Obs = .N,
  Failures = sum(as.numeric(failure), na.rm = TRUE),
  Failure_Rate_pct = (sum(as.numeric(failure), na.rm = TRUE) / .N) * 100
), by = capacity_tb][order(capacity_tb)]

```

```

# Dynamic retrieval of the list of capacities for the subtitle (e.g. '12TB, 14TB, 16TB,
↳ 20TB')
cap_list_str <- paste0(cap_stats$capacity_tb, "TB", collapse = ", ")
sub_title_f6 <- sprintf("Comparison across detected operational drive sizes (%s)",
↳ cap_list_str)

p6 <- ggplot(cap_stats, aes(x = factor(capacity_tb), y = Failure_Rate_pct)) +
  geom_col(fill = "#2ecc71", width = 0.5) +
  geom_text(aes(label = sprintf("%.4f%%", Failure_Rate_pct)), vjust = -0.5, size = 3.8,
↳ fontface = "bold") +
  scale_y_continuous(expand = expansion(mult = c(0, 0.25)), labels = function(x)
↳ paste0(sprintf("%.4f", x), "%")) +
  labs(
    title = "Hard Drive Failure Rate by Capacity (TB)",
    subtitle = sub_title_f6,
    x = "Storage Capacity (TB)",
    y = "Daily Failure Rate (%)"
  )

ggsave(file.path(PATH_FIGS, "fig6_failure_rate_by_capacity.png"), plot = p6, width =
↳ max(8, nrow(cap_stats) * 1.2), height = 5, dpi = 300)
rm(dt_cap, cap_stats)

cat("\n[SUCCESS] All 6 EDA visualisations generated and saved to 'outputs/figures/'\n")
}

```

```

gc(verbose = FALSE)

```

4.2.1 Class Imbalance and Model-Level Vulnerabilities




Figure 1: Class Imbalance between Operational Logs and 14-Day Failure Horizon

As demonstrated in Figure 1, the dataset exhibits extreme target class imbalance. Out of 13,172,916 daily observations in the filtered Top 5 fleet, healthy operational days account for 13,167,669 records (99.960%), whereas positive failure window records ($Y_{it} = 1$) constitute only 5,247 observations (0.040%). This ratio ($\approx 2,509 : 1$) highlights that standard accuracy metrics would yield misleadingly high performance (99.96%) for a naive baseline classifier predicting all drives as healthy. Consequently, modeling evaluation must focus strictly on Precision-Recall AUC (PR-AUC) and operational SRE cost matrices rather than accuracy or ROC-AUC alone.




Figure 2: Daily Failure Rate Heterogeneity across Hard Drive Models

Figure 2 illustrates significant hardware failure rate heterogeneity across the top drive models in the filtered fleet:

- **Seagate ST12000NM0008 (12 TB)** exhibits the highest daily failure rate at 0.0082%, making it nearly $3\times$ more prone to failure than Toshiba models.
- **Toshiba MG07ACA14TA (14 TB)** and **Toshiba MG08ACA16TA (16 TB)** demonstrate moderate daily failure rates of 0.0029% and 0.0025%, respectively.
- **Seagate ST16000NM001G (16 TB)** shows a low daily failure rate of 0.0019%.
- **Western Digital WDC WUH721816ALE6L4 (16 TB)** displays the lowest failure rate in the fleet at 0.0008%, proving exceptional structural stability.

4.2.2 SMART Health Distributions and Collinearity Analysis




Figure 3: Log-Transformed SMART Feature Distributions: Failing vs Healthy Drives

The comparative distribution analysis in Figure 3 confirms distinct physical degradation behaviors prior to drive failure:

- **SMART 5 (Reallocated Sectors Count):** Healthy drives remain strictly centered at zero ($\log1p = 0$), whereas failing drives display a substantially elevated interquartile range (median $\log1p \approx 2.1$), signaling physical disk surface degradation.

- **SMART 187 (Reported Uncorrectable Errors):** Healthy drives show virtually no uncorrectable read errors, while failing drives exhibit strong median spikes ($\log_{10} p \approx 2.5$), serving as an immediate physical failure indicator.
- **SMART 197 (Current Pending Sector Count):** Displays a clear upward shift for failing drives, representing unstable sectors awaiting remapping.
- **SMART 194 (Temperature °C):** Thermal distributions remain nearly identical between failing (median 31.5°C) and healthy drives (median 33.0°C), demonstrating that ambient operating temperature is not a direct primary trigger for acute hardware failure in this datacenter environment.




Figure 4: Spearman Correlation Matrix among Key SMART Telemetry Features

Figure 4 presents the Spearman rank correlation matrix across key raw SMART metrics and their corresponding 7-day temporal deltas (Δ_7). Three important structural collinearity patterns emerge:

1. **Moderate Collinearity ($\rho = 0.49$) between SMART 197 (Pending Sectors) and SMART 198 (Uncorrectable Sectors):** These two indicators measure closely linked sector-level degradation dynamics. While they capture distinct stages of sector instability, their moderate positive correlation indicates shared predictive variance, validating feature selection or tree-based regularisation to prevent redundant splitting.
2. **Physical Degradation Coupling ($\rho = 0.33$) between SMART 5 and SMART 187:** This reflects the physical transition process where uncorrectable read errors (`smart_187_raw`) lead to firmware reallocations of bad sectors (`smart_5_raw`). A similar secondary coupling is observed between sector errors and their temporal rates of change, such as `smart_197_raw` and its 7-day delta `smart_197_delta7` ($\rho = 0.36$).
3. **Orthogonality of Operational Controls ($\rho \approx -0.00$ to -0.11):** Baseline operational attributes like temperature (`smart_194_raw`) and power-on time (`smart_9_raw`) display near-zero linear redun-

dancy with acute failure indicators, confirming that thermal metrics and drive age operate independently of sector-level hardware degradation.

4.2.3 Temporal Dynamics and Capacity Risk Profiles




Figure 5: Daily Failure Rate Dynamics and LOESS Trend Line across Q1 2024

Figure 5 plots the daily failure rate time series across Q1 2024. The daily failure rate fluctuates within a tight operational band between 0.0010% and 0.0085%. The non-parametric LOESS smoothing line reveals a steady, stable global baseline centered around 0.0030%–0.0040%, with minor transient spikes in early February (0.0084%). This temporal stability confirms the absence of datacenter-wide systemic shocks (e.g., power outages or environmental failures), validating the dataset for time-based train-validation splitting.




Figure 6: Daily Failure Rate Comparison across Operational Storage Capacities

Figure 6 breaks down daily failure rates by drive capacity (12 TB, 14 TB, and 16 TB):

- **12 TB Drives:** Highest failure rate (0.0082%), driven entirely by the aging **ST12000NM0008** model fleet.
- **14 TB Drives:** Intermediate failure rate (0.0029%), corresponding to the **TOSHIBA MG07ACA14TA** workload.
- **16 TB Drives:** Lowest aggregate failure rate (0.0019%), reflecting newer enterprise architecture across Toshiba, Seagate, and WDC 16 TB models.

4.2.4 Predictive Window and Feature Engineering Design

The exploratory data analysis and the subsequent modeling pipeline both use the same 14-day target horizon (`PREDICTION_WINDOW_DAYS <- 14, target_14d`), ensuring consistency between EDA-level failure distributions and the label used for training and evaluation. Separately, feature engineering relies on shorter, 7-day rolling deltas (Δ_7) computed on critical S.M.A.R.T. indicators (`smart_5_delta7`, `smart_187_delta7`, `smart_197_delta7`). This shorter look-back window was chosen because physical degradation mechanisms, particularly pending sector accelerations, tend to manifest over short time spans, making a 7-day rate-of-change a higher-frequency, less noisy predictor than a 14-day one, while remaining computationally tractable within the 6 GB virtual machine constraint. This distinction between the target horizon (H) and the feature look-back window (Δ_7) is formalized in Section 5.1. But, in short: $H = 14$ days defines what the model predicts (the label), while Δ_7 defines how far back an input feature looks ; the two are independent design choices.

5 Machine Learning and Predictive Modelling in an SRE Context

This chapter presents the modelling strategy used to forecast imminent hard drive failures within a 14-day proactive maintenance window. The approach integrates temporal label construction, strict leakage-free train/test separation, class-imbalance mitigation, and reliability-oriented evaluation metrics. The modelling pipeline is designed to align with SRE operational constraints, where false negatives carry significant SLA penalties and proactive replacement reduces toil and protects error budgets.

5.1 Construction of the Predictive Failure Target ($H = 14$ Days)

To train a proactive failure model, binary target labels must reflect future failure events within an $H = 14$ -day projection window. For each drive i observed at date t , the prospective failure target Y_{it} is defined as:

$$Y_{it} = \max_{k \in [1, 14]} (\text{failure}_{i, t+k})$$

Where $\text{failure}_{i, t+k} \in \{0, 1\}$ represents the ground-truth terminal status of drive i at time $t + k$. This formulation transforms a point-in-time anomaly detection task into a time-series binary classification problem, giving data center administrators a 14-day operational lead time for drive replacement.

Methodological distinction: Forecast horizon (H) vs. Analysis window (Δ_7) It is essential to distinguish between the role of the forecast horizon (H) and that of the feature aggregation window (Δ_7). These two parameters serve distinct purposes and are defined on either side of the prediction time t :

- **The target horizon ($H = 14$ days) - The future trajectory ($t \rightarrow t + 14$):** This determines the event that the model seeks to predict. A positive label ($Y_{i,t} = 1$) indicates that a failure will occur within the window $[t, t + 14]$. The value of this horizon is dictated by SRE operational constraints: it guarantees the minimum time required to schedule the hot drain of a node, rebuild parity and isolate the disk without impacting service availability.
- **The feature window ($\Delta_7 = 7$ days) - Recent history ($t - 7 \rightarrow t$):** This defines the temporal depth used to capture the dynamics of hardware degradation. Rather than relying on a snapshot SMART measurement at time t , the Δ_7 window calculates time derivatives and accumulation rates (e.g. the rate of increase in reallocated sectors `smart_5_delta7` or uncorrelated read errors `smart_187_delta7` over the last 7 days).

In summary, Δ_7 measures the rate of past degradation to feed into the model, whilst H defines the lead time required by the business. The historical aggregation (Δ_7) and the lead time (H) are conceptually independent: varying the observation window in no way alters the intervention window required by the infrastructure.

```
#We return only a list at the end of `local()` containing only the key figures we need
pipeline_metrics <- local({
  # =====
  # Script: 02_model_training.R
  # Project: Predictive Hard Drive Failure Modelling for SRE Operations
  # Description: Strict temporal train/test split, model training (Baseline & XGBoost)
  # =====

  # 0. Load Libraries
  required_packages <- c("data.table", "tidyverse", "lubridate", "xgboost", "pROC")

  for (pkg in required_packages) {
```



```

if (!require(pkg, character.only = TRUE)) {
  install.packages(pkg)
  library(pkg, character.only = TRUE)
}
}

options(scipen = 999)
gc()

# =====
# 1. Global Configuration & Paths
# =====
DATASET_TAG    <- "Q1_2024"
PATH_PROCESSED <- "data/processed/"
PATH_MODELS    <- "outputs/models/"
PATH_METRICS   <- "outputs/metrics/"

if (!dir.exists(PATH_MODELS)) dir.create(PATH_MODELS, recursive = TRUE)

INPUT_FILEPATH <- file.path(PATH_PROCESSED, sprintf("dt_processed_%s.rds",
↳ tolower(DATASET_TAG)))

# =====
# 2. Loading Processed Dataset
# =====
cat(sprintf("Loading processed dataset from %s...\n", INPUT_FILEPATH))
dt <- readRDS(INPUT_FILEPATH)

dt[, date := as.IDate(date)]

# =====
# 3. Strict Temporal Train / Test Split (Prevention of Temporal Leakage)
# =====
cutoff_date <- as.IDate("2024-03-01")

cat(sprintf("\n--- Performing Strict Temporal Split at %s ---\n", cutoff_date))

train_dt <- dt[date < cutoff_date]
test_dt  <- dt[date >= cutoff_date]

cat(sprintf("Train Set: %s rows | Failure Window (14d) Count: %s\n",
  format(nrow(train_dt), big.mark = " "), sum(train_dt$target_14d)))
cat(sprintf("Test Set : %s rows | Failure Window (14d) Count: %s\n",
  format(nrow(test_dt), big.mark = " "), sum(test_dt$target_14d)))

rm(dt)
gc()

# =====
# 4. Feature Selection & Preprocessing
# =====
cat("\n--- Preparing features & handling missing values ---\n")

train_dt[, capacity_tb := capacity_bytes / 1e12]

```

```

test_dt[, capacity_tb := capacity_bytes / 1e12]

feature_cols <- c(
  "capacity_tb", "smart_5_raw", "smart_9_raw", "smart_187_raw",
  "smart_188_raw", "smart_194_raw", "smart_197_raw", "smart_198_raw",
  "smart_5_delta7", "smart_187_delta7", "smart_197_delta7"
)

target_col <- "target_14d"

# Impute NA values with 0 directly in-place
for (j in feature_cols) {
  set(train_dt, which(is.na(train_dt[[j]])), j, 0)
  set(test_dt, which(is.na(test_dt[[j]])), j, 0)
}

# =====
# 5. Baseline Model: Logistic Regression (Anti-Collinearity & RAM-Optimised)
# =====
cat("\n--- Training Model 1: Logistic Regression (Baseline) ---\n")

set.seed(42)
pos_indices <- train_dt[target_14d == 1, which = TRUE]
neg_indices <- train_dt[target_14d == 0, which = TRUE]

# Subsample 20x negative instances relative to positive count (~80k rows)
sampled_neg_indices <- sample(neg_indices, size = length(pos_indices) * 20)
lr_train_dt <- train_dt[c(pos_indices, sampled_neg_indices)]

# 5.1 Automatic Detection & Dropping of Constant or Collinear Features for GLM
glm_features <- c()
for (col in feature_cols) {
  vals <- log1p(pmax(0, lr_train_dt[[col]]))
  if (sd(vals) > 0.0001) {
    glm_features <- c(glm_features, col)
  }
}

# Drop collinear columns with correlation > 0.98 in training slice
if (length(glm_features) > 1) {
  # Convert to matrix and clamp negative values without flattening matrix dimensions
  mat <- as.matrix(lr_train_dt[, ..glm_features])
  mat[mat < 0] <- 0

  cor_mat <- cor(log1p(mat))
  cor_mat[upper.tri(cor_mat, diag = TRUE)] <- 0

  collinear_cols <- glm_features[apply(abs(cor_mat) > 0.98, 2, any)]
  if (length(collinear_cols) > 0) {
    cat(sprintf("Dropping highly collinear features for GLM: %s\n", paste(collinear_cols,
      ↪ collapse = ", ")))
    glm_features <- setdiff(glm_features, collinear_cols)
  }
}

```

```

# 5.2 Scale Subsampled Train Set for GLM
scaled_train_dt <- copy(lr_train_dt[, ..glm_features])
scale_means <- list()
scale_sds <- list()

for (col in glm_features) {
  raw_vals <- log1p(pmax(0, scaled_train_dt[[col]]))
  m_val <- mean(raw_vals)
  s_val <- sd(raw_vals)
  if (s_val == 0) s_val <- 1

  scale_means[[col]] <- m_val
  scale_sds[[col]] <- s_val

  scaled_train_dt[, (col) := (raw_vals - m_val) / s_val]
}
scaled_train_dt[, (target_col) := lr_train_dt[[target_col]]]

# 5.3 Fit Logistic Regression
lr_formula <- as.formula(paste(target_col, "~", paste(glm_features, collapse = " + ")))

lr_model <- glm(
  formula = lr_formula,
  data = scaled_train_dt,
  family = binomial(link = "logit")
)

cat("Logistic Regression fitted successfully. Predicting on Test Set...\n")

# 5.4 Scale Test Set on-the-fly and Predict
scaled_test_dt <- data.table(matrix(0, nrow = nrow(test_dt), ncol =
  ↪ length(glm_features)))
colnames(scaled_test_dt) <- glm_features

for (col in glm_features) {
  t_vals <- log1p(pmax(0, test_dt[[col]]))
  scaled_test_dt[, (col) := (t_vals - scale_means[[col]]) / scale_sds[[col]]]
}

test_dt[, pred_lr := predict(lr_model, newdata = scaled_test_dt, type = "response")]
cat("Baseline predictions generated successfully.\n")

# Free memory from intermediate objects
rm(lr_train_dt, scaled_train_dt, scaled_test_dt)
gc()

# =====
# 6. Advanced Model: XGBoost Classifier (Full Train Set)
# =====
cat("\n--- Training Model 2: XGBoost Classifier (Full 8.5M Train Set) ---\n")

# Prepare DMatrix objects
dtrain <- xgb.DMatrix(

```

```

    data = as.matrix(train_dt[, ..feature_cols]),
    label = train_dt[[target_col]]
)

dtest <- xgb.DMatrix(
  data = as.matrix(test_dt[, ..feature_cols]),
  label = test_dt[[target_col]]
)

# Calculate class imbalance ratio for scale_pos_weight
neg_count <- sum(train_dt[[target_col]] == 0)
pos_count <- sum(train_dt[[target_col]] == 1)
scale_pos_weight_val <- neg_count / pos_count

xgb_params <- list(
  booster      = "gbtree",
  objective     = "binary:logistic",
  eval_metric   = "auc",
  max_depth    = 6,
  eta          = 0.1,
  scale_pos_weight = scale_pos_weight_val,
  nthread      = 2
)

cat(sprintf("Class Imbalance Ratio (scale_pos_weight): %.2f\n", scale_pos_weight_val))

set.seed(42)
xgb_model <- xgb.train(
  params      = xgb_params,
  data        = dtrain,
  nrounds     = 100,
  watchlist   = list(train = dtrain, eval = dtest),
  early_stopping_rounds = 10,
  print_every_n = 20,
  verbose     = 0
)

# Predict probabilities on FULL Test Set
test_dt[, pred_xgb := predict(xgb_model, dtest)]

# modification of script02 to extract the information needed for Markdown
n_train_val <- nrow(train_dt)
n_test_val  <- nrow(test_dt)

rm(dtrain, dtest, train_dt)
gc()

# Feature Importance & Metrics Export
importance_matrix <- xgb.importance(model = xgb_model)

engineered_cols <- c("smart_5_delta7", "smart_187_delta7", "smart_197_delta7")
gain_engineered <- sum(importance_matrix[Feature %in% engineered_cols, Gain])
gain_total      <- sum(importance_matrix$Gain)

```

```

feature_importance_engineered <- gain_engineered / gain_total

if (!dir.exists(PATH_METRICS)) dir.create(PATH_METRICS, recursive = TRUE) # Ensure that
  ↳ the metrics directory exists
write.csv(
  data.frame(feature_importance_engineered = feature_importance_engineered),
  file.path(PATH_METRICS, "feature_importance_summary.csv"),
  row.names = FALSE
)
# Explicit memory clearance
rm(importance_matrix, gain_engineered, gain_total, engineered_cols)
gc(verbose = FALSE)

# =====
# 7. Save Models and Test Predictions
# =====
cat("\n--- Saving Trained Models & Test Predictions ---\n")

saveRDS(lr_model, file = file.path(PATH_MODELS, "model_logistic_regression.rds"))
xgb.save(xgb_model, file.path(PATH_MODELS, "model_xgboost.model"))

output_preds_path <- file.path(PATH_PROCESSED, sprintf("test_predictions_%s.rds",
  ↳ tolower(DATASET_TAG)))
saveRDS(test_dt[, .(serial_number, date, model, target_14d, pred_lr, pred_xgb)], file =
  ↳ output_preds_path)

cat(sprintf("Predictions successfully exported to %s\n", output_preds_path))
cat("Phase 3 Training pipeline completed successfully!\n")
gc()

# Extraction sécurisée du log d'évaluation et de l'early stopping pour le rapport
eval_log <- attr(xgb_model, "evaluation_log")
b_iter <- attr(xgb_model, "early_stop")$best_iteration

b_score <- NA_real_

if (!is.null(eval_log) && nrow(eval_log) > 0) {
  # Detection dynamique de la colonne de score d'évaluation (ignore le nom exact)
  score_col <- grep("eval|test", names(eval_log), value = TRUE)[1]

  if (is.null(b_iter) || is.na(b_iter)) {
    b_iter <- eval_log[[1]][which.max(eval_log[[score_col]])]
  }

  if (!is.null(score_col) && !is.na(score_col)) {
    b_score <- eval_log[eval_log[[1]] == b_iter, get(score_col)]
  }
}

# Retour de la liste finale pour pipeline_metrics
list(
  n_train      = n_train_val,

```

```

n_test      = n_test_val,
scale_pos_weight = scale_pos_weight_val,
best_iter   = as.integer(b_iter),
best_score  = as.numeric(b_score)
)
})
# 3. SAUVEGARDE RDS (ajoutée juste ici, hors du local)
if (!dir.exists("outputs/models")) dir.create("outputs/models", recursive = TRUE)
saveRDS(pipeline_metrics, "outputs/models/pipeline_metrics.rds")
invisible(gc(verbose = FALSE))

```

Technical Implementation & Pipeline Optimization Note:

To seamlessly integrate the core training logic (02_model_training.R) into the dynamic reporting workflow, three key adjustments were implemented:

- Object Structure & C++ API Extraction:** The native xgboost model object (xgb.Booster) is a C++ external pointer structure (externalptr)¹, meaning internal training attributes are not exposed via standard R list fields (e.g., xgb_model\$best_iteration). Instead, the optimal evaluation epoch and metrics were dynamically extracted using R internal attributes (attr(xgb_model, “evaluation_log”) and attr(xgb_model, “early_stop”)), guaranteeing accurate metrics retrieval regardless of minor package version variations.
- Isolated Execution Environment:** Running the training script inside a local({ ... }) block prevents memory leakage and variable collisions with the global environment, ensuring clean environment state management during compilation.
- Artifact Persistence & Memory Protection:** To prevent out-of-memory (OOM) crashes during Knitr compilation on large matrix operations (~8.5M rows), key pipeline outputs are compiled into a lightweight pipeline_metrics object and persisted to disk (pipeline_metrics.rds). This decouples heavy computations from report rendering, allowing deterministic and fast document generation.

5.1.1 Strict Temporal Train/Test Splitting Protocol

In time-series operational forecasting, partitioning was performed along a hard temporal boundary set at **March 1, 2024** to prevent temporal leakage.

Parameter / Configuration	Baseline (Logistic Regression)	Advanced (XGBoost Classifier)
Training Sample Size	Subsampled 20x Negative (~80k rows)	Full Matrix (8 537 677 rows)
Test Sample Size	Full Test Partition	Full Matrix (4 635 239 rows)
Feature Set Strategy	Standardized log1p (Filtered p < 0.98)	Full Raw + 7-Day Rolling Deltas (D7)
Imbalance Treatment	Class Subsampling	Cost Weighting (scale_pos_weight = 2157.16)
Tree Hyperparameters	N/A	max_depth = 6, eta = 0.1, nrounds = 100
Best XGBoost Iteration	N/A	Iteration 24 (Eval AUC = 0.8364)

5.1.2 Feature Engineering: The 7-Day SMART Delta (Δ_7)

To capture short-term acceleration in drive degradation independently of the 14-day target horizon, a 7-day rolling window delta (Δ_7) was engineered as the core dynamic predictor:

¹Chen, T., & Guestrin, C. (2016). *XGBoost: A Scalable Tree Boosting System*. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. XGBoost R Package C++ API Reference: <https://xgboost.readthedocs.io/en/stable/R-package/>

$$\Delta_7 SMART_{i,t} = SMART_{i,t} - SMART_{i,t-7}$$

Rapid positive accelerations in $\Delta_7 SMART_{197}$ or $\Delta_7 SMART_{187}$ provide actionable short-term failure signals, giving SRE teams a 14-day lead time to queue drive replacements.

5.2 Model Evaluation: Reliability Metrics, ROC/PR Curves and Temporal Generalisation

The evaluation module assesses the predictive performance of the trained models using reliability-oriented metrics aligned with SRE operational constraints. Because hard drive failures represent less than 0.04% of daily observations, classical accuracy is not meaningful; instead, the evaluation focuses on AUC-ROC, Precision-Recall curves, and PR-AUC, which quantify the model's ability to detect rare failure events without generating excessive false positives.

The script loads the temporally separated test set and computes ROC curves for both the baseline logistic regression and the XGBoost classifier. The ROC comparison highlights the superior ability of tree-based models to capture non-linear degradation patterns in S.M.A.R.T. telemetry. However, given the extreme class imbalance, ROC curves alone are insufficient. The module therefore computes a Precision-Recall curve, which provides a more realistic view of model performance under operational conditions where recall is critical and precision must remain acceptable.

From an SRE perspective, these metrics directly reflect the asymmetric cost structure of hardware failures: a false negative triggers emergency RAID rebuilds and threatens SLO compliance, whereas a false positive incurs only marginal hardware replacement cost. The evaluation module therefore prioritises high recall within acceptable precision ranges, ensuring that predictions remain actionable for proactive maintenance workflows.

The script exports a summary table consolidating AUC-ROC, PR-AUC, and failure counts, enabling transparent reporting and reproducibility within the thesis.

To evaluate the predictive performance of the models under real-world SRE operational conditions, all evaluations were strictly performed on the out-of-time temporal test split (March 2024, comprising 4,635,239 disk-day observations).

```
local({
# =====
# Script: 03_model_evaluation.R
# Project: Predictive Hard Drive Failure Modelling for SRE Operations
# Description: Model Evaluation (AUC-ROC, Precision-Recall, SRE Cost Curves)
# =====

# 0. Load Libraries
required_packages <- c("data.table", "tidyverse", "pROC", "PRROC", "scales", "gridExtra")

for (pkg in required_packages) {
  if (!require(pkg, character.only = TRUE)) {
    install.packages(pkg)
    library(pkg, character.only = TRUE)
  }
}

theme_set(theme_minimal(base_size = 12))

PATH_PROCESSED <- "data/processed/"
```

```

PATH_FIGS      <- "outputs/figures/"
PATH_METRICS   <- "outputs/metrics/"

if (!dir.exists(PATH_METRICS)) dir.create(PATH_METRICS, recursive = TRUE)

INPUT_PREDS <- file.path(PATH_PROCESSED, "test_predictions_q1_2024.rds")

cat("Loading test set predictions...\n")
preds_dt <- readRDS(INPUT_PREDS)

# =====
# 1. Compute ROC Curves & AUC Metric
# =====
cat("\n--- Calculating ROC & PR Metrics ---\n")

roc_lr <- roc(preds_dt$target_14d, preds_dt$pred_lr, quiet = TRUE)
roc_xgb <- roc(preds_dt$target_14d, preds_dt$pred_xgb, quiet = TRUE)

auc_lr_val <- auc(roc_lr)
auc_xgb_val <- auc(roc_xgb)

cat(sprintf("Baseline (Logistic Regression) AUC-ROC: %.4f\n", auc_lr_val))
cat(sprintf("Advanced (XGBoost) AUC-ROC: %.4f\n", auc_xgb_val))

# =====
# 2. Plotting ROC Curves Comparison
# =====
df_roc_lr <- data.frame(FPR = 1 - roc_lr$specificities, TPR = roc_lr$sensitivities,
  ↪ Model = sprintf("Logistic Regression (AUC = %.3f)", auc_lr_val))
df_roc_xgb <- data.frame(FPR = 1 - roc_xgb$specificities, TPR = roc_xgb$sensitivities,
  ↪ Model = sprintf("XGBoost (AUC = %.3f)", auc_xgb_val))

df_roc <- rbind(df_roc_lr, df_roc_xgb)

p_roc <- ggplot(df_roc, aes(x = FPR, y = TPR, color = Model)) +
  geom_line(size = 1.2) +
  geom_abline(slope = 1, intercept = 0, linetype = "dashed", color = "gray50") +
  scale_color_manual(values = c("#e74c3c", "#2980b9")) +
  labs(
    title = "ROC Curves Comparison on Temporal Test Set (March 2024)",
    x = "False Positive Rate (1 - Specificity)",
    y = "True Positive Rate (Sensitivity / Recall)",
    color = "Model"
  ) +
  theme(legend.position = "bottom")

ggsave(file.path(PATH_FIGS, "fig7_roc_curves_comparison.png"), plot = p_roc, width = 8,
  ↪ height = 6, dpi = 300)

# =====
# 3. Precision-Recall Curve (Critical for Extreme Class Imbalance)
# =====
pr_lr <- pr.curve(
  scores.class0 = preds_dt[target_14d == 1, pred_lr],

```



```

    scores.class1 = preds_dt[target_14d == 0, pred_lr],
    curve = TRUE
  )

pr_xgb <- pr.curve(
  scores.class0 = preds_dt[target_14d == 1, pred_xgb],
  scores.class1 = preds_dt[target_14d == 0, pred_xgb],
  curve = TRUE
)

df_pr <- data.frame(Recall = pr_xgb$curve[, 1], Precision = pr_xgb$curve[, 2])

p_pr <- ggplot(df_pr, aes(x = Recall, y = Precision)) +
  geom_line(color = "#27ae60", size = 1.2) +
  labs(
    title = sprintf("XGBoost Precision-Recall Curve (PR-AUC = %.4f)",
      ↪ pr_xgb$auc.integral),
    subtitle = "Essential evaluation metric given the 0.04% positive class imbalance",
    x = "Recall (Sensitivity)",
    y = "Precision"
  )

ggsave(file.path(PATH_FIGS, "fig8_precision_recall_curve.png"), plot = p_pr, width = 8,
  ↪ height = 6, dpi = 300)

# =====
# 4. Export Summary Table for Thesis Report
# =====
summary_table <- data.table(
  Model = c("Logistic Regression (Baseline)", "XGBoost Classifier"),
  AUC_ROC = c(auc_lr_val, auc_xgb_val),
  PR_AUC = c(NA, pr_xgb$auc.integral),
  Test_Observations = nrow(preds_dt),
  Test_Failures = sum(preds_dt$target_14d)
)

write.csv(summary_table, file.path(PATH_METRICS, "model_performance_summary.csv"),
  ↪ row.names = FALSE)

cat("\nSummary of results exported to 'outputs/metrics/model_performance_summary.csv'\n")
cat("Phase 4 Model Evaluation completed successfully!\n")
})
gc(verbose = FALSE)

```

5.2.1 Discriminative Power (ROC Curves Analysis)

The Receiver Operating Characteristic (ROC) curve evaluates the trade-off between the True Positive Rate (Recall) and the False Positive Rate across all decision thresholds.




Figure 7: ROC Curves Comparison on Temporal Test Set (March 2024)

Both models demonstrate strong discrimination capabilities over a random baseline:

- **Baseline (Logistic Regression):** Achieves an AUC-ROC of 0.818, showing that log-transformed SMART features combined with subsampling provide a solid linear baseline.
- **XGBoost Classifier:** Reaches an AUC-ROC of 0.836, outperforming the baseline by capturing non-linear feature interactions and non-monotonic failure degradation patterns.

5.2.2 Precision-Recall Dynamics in Extreme Class Imbalance

While AUC-ROC provides a general measure of ranking capability, it can present an overly optimistic view in datasets with extreme class imbalance ($\sim 0.04\%$ failure prevalence). Because the total number of true negatives is massive ($\sim 4.6\text{M}$ rows), large shifts in false positives produce negligible changes in the False Positive Rate ($FPR = \frac{FP}{FP+TN}$). To assess operational reliability for Site Reliability Engineering (SRE) teams, the Precision-Recall (PR) Curve provides a far more informative evaluation.




Figure 8: XGBoost Precision-Recall Curve under 0.04% Class Imbalance

Key Operational Takeaways from the PR Curve (PR-AUC = 0.0516):

- **Baseline Expectation:** In a dataset with 0.04% positive prevalence, a random classifier yields a Precision equal to the prevalence (0.0004). An Area Under the Precision-Recall Curve (PR-AUC) of 0.0516 represents a $\sim 129\times$ lift over random guessing.
- **High-Precision Regime (Low Recall):** At very low recall thresholds ($< 5\%$), the model achieves high precision ($> 75\%$), making it suitable for fully automated, high-confidence maintenance actions (e.g., immediate drive evacuation).
- **Operational Alerting Regime (20% - 30% Recall):** In the operational window where SRE teams capture 20% to 30% of impending drive failures, precision stabilizes between 10% and 20%. This indicates that for every 1 true failure detected, approximately 4 to 9 false alarms are raised.

5.2.3 Performance Summary

The consolidated metrics across both models are summarized below:

Table 5: Model Performance Summary on Out-of-Time Test Set (March 2024)

Model	Test Observations	Test Failures	AUC-ROC	PR-AUC
Logistic Regression (Baseline)	4 635 239	1 291	0.8176	N/A
XGBoost Classifier	4 635 239	1 291	0.8364	0.0516

A comparative analysis of the metrics confirms the overall superiority of the XGBoost approach over the baseline model. By achieving an AUC-ROC of 0.8364 compared with 0.8176 for logistic regression, XGBoost

demonstrates better discriminatory power. On an operational scale of over 4.63 million observations, this 1.88 per cent improvement results in a significantly more accurate distinction between healthy disks and those in the process of failing, thereby reducing false alarms whilst maximising the detection of actual failures.

This performance can be attributed to the intrinsic ability of decision trees to capture complex dynamics that logistic regression cannot model. Whereas the baseline model relies on a strict linear relationship applied to transformed variables ($\log(1+x)$), XGBoost automatically learns the non-linear interactions between several SMART metrics, as well as the non-monotonic deterioration patterns provided by the 7-day moving deltas (Δ_7).

Finally, the key difference lies in the handling of class imbalance. To converge, logistic regression required a selective subsample of the negative class (~80k rows), thereby sacrificing a large portion of the data's variability. Conversely, XGBoost was trained directly on the full matrix of 8.5 million rows by adjusting the cost via the parameter `scale_pos_weight = 2157.16`. This full utilisation of the signal ensures greater model robustness in the face of the diversity of behaviour exhibited by healthy disks in production.

SRE Decision Framework Note: In data center operations, a Precision of ~15% is highly actionable. The cost of a proactive, non-disruptive hot-swap drive replacement is orders of magnitude lower than an ungraceful drive crash leading to RAID rebuilding, network saturation, or data loss.

5.3 SRE Cost-Benefit Analysis: Operational Economics and Threshold Optimisation

This module quantifies the operational and economic impact of predictive maintenance decisions by modelling the asymmetric costs associated with false negatives, false positives, and true positives. Using calibrated XGBoost probabilities, the script evaluates a wide range of decision thresholds and computes the total maintenance cost under each scenario. This enables the identification of the optimal threshold that minimises global SRE cost.

Two operational scenarios are analysed. The first represents a standard hardware replacement environment, where the cost of an unpredicted failure (FN) is moderate. The second models an enterprise-grade environment with strict SLA requirements, where FN costs are significantly higher due to RAID rebuild overhead, degraded latency, and potential cascading failures. The sensitivity analysis demonstrates how higher FN penalties shift the optimal threshold toward more aggressive detection, increasing recall while maintaining acceptable precision.

From an SRE perspective, this module bridges the gap between statistical performance and real-world operational impact. It quantifies how proactive prediction reduces emergency interventions, protects error budgets, and eliminates toil. The cost curves illustrate the trade-off between proactive and reactive strategies, showing that even modest recall improvements can yield substantial economic savings in environments with high SLA penalties.

The script exports a consolidated metrics table summarising optimal thresholds, cost savings, and confusion-matrix components, enabling integration into the thesis and supporting decision-making for SRE teams.

```
local({
# =====
# Script: 04_sre_cost_analysis.R
# Project: Predictive Hard Drive Failure Modelling for SRE Operations
# Description: SRE Cost-Benefit Analysis & Economic Sensitivity Analysis
# =====

# 0. Load Libraries
required_packages <- c("data.table", "tidyverse", "scales", "gridExtra", "xgboost")
```

```

for (pkg in required_packages) {
  if (!require(pkg, character.only = TRUE)) {
    install.packages(pkg)
    library(pkg, character.only = TRUE)
  }
}

theme_set(theme_minimal(base_size = 12))

PATH_PROCESSED <- "data/processed/"
PATH_FIGS      <- "outputs/figures/"
PATH_METRICS   <- "outputs/metrics/"
PATH_MODELS    <- "outputs/models/"

if (!dir.exists(PATH_METRICS)) dir.create(PATH_METRICS, recursive = TRUE)

INPUT_PREDS <- file.path(PATH_PROCESSED, "test_predictions_q1_2024.rds")

cat("Loading test set predictions...\n")
preds_dt <- readRDS(INPUT_PREDS)

# =====
# 1. Cost Optimization Function
# =====
run_cost_optimization <- function(preds_dt, cost_fn, cost_fp, cost_tp, scenario_label) {

  # 1. Calculation of the calibrated probability (XGBoost scale correction)
  POS_WEIGHT <- (1 - 0.0004) / 0.0004
  y_prob <- preds_dt$pred_xgb / (preds_dt$pred_xgb + ((1 - preds_dt$pred_xgb) *
    ↪ POS_WEIGHT))

  # 2. Threshold scanning adapted to very low calibrated probabilities
  thresholds <- unique(c(
    seq(0.000001, 0.0001, length.out = 100),
    seq(0.0001, 0.01, length.out = 100),
    seq(0.01, 0.5, length.out = 50)
  ))

  y_true <- preds_dt$target_14d
  n_total <- length(y_true)
  n_failures <- sum(y_true)

  cost_results <- vector("list", length(thresholds))

  for (i in seq_along(thresholds)) {
    t <- thresholds[i]
    pred_pos <- y_prob >= t

    tp <- sum(pred_pos & y_true == 1)
    fp <- sum(pred_pos & y_true == 0)
    fn <- sum(!pred_pos & y_true == 1)
    tn <- sum(!pred_pos & y_true == 0)
  }
}

```

```

total_cost <- (tp * cost_tp) + (fp * cost_fp) + (fn * cost_fn)
cost_per_drive <- total_cost / n_total

precision <- ifelse((tp + fp) > 0, tp / (tp + fp), 0)
recall <- ifelse((tp + fn) > 0, tp / (tp + fn), 0)
f1 <- ifelse((precision + recall) > 0, 2 * (precision * recall) / (precision
↪ + recall), 0)

cost_results[[i]] <- data.table(
  Scenario = scenario_label,
  threshold = t,
  TP = tp, FP = fp, FN = fn, TN = tn,
  total_cost = total_cost,
  cost_per_drive = cost_per_drive,
  precision = precision,
  recall = recall,
  f1 = f1
)
}

cost_df <- rbindlist(cost_results)
reactive_cost <- n_failures * cost_fn
optimal_row <- cost_df[which.min(total_cost)]

return(list(
  cost_df = cost_df,
  optimal_row = optimal_row,
  reactive_cost = reactive_cost,
  cost_fn = cost_fn,
  cost_fp = cost_fp,
  cost_tp = cost_tp
))
}

# =====
# 2. Execution of Both Scenarios (Sensitivity Analysis)
# =====
cat("\n--- Running Pass 1: Standard Operational Model (C_FN = $100) ---\n")
# JUSTIFICATION: Baseline hardware replacement cost without SLA impact penalties.
# At C_FN=$100, precision threshold needed is 14.29%. High FP penalty dominates.
res_s1 <- run_cost_optimization(preds_dt, cost_fn = 100, cost_fp = 15, cost_tp = 10,
                               scenario_label = "Scenario A: Standard ($100 FN)")

cat("\n--- Running Pass 2: Enterprise SLA/RAID Model (C_FN = $500) ---\n")
# JUSTIFICATION: Real-world SRE environments where unplanned downtime, RAID rebuilds,
# and potential cascading failures carry heavy SLA penalties ($500+).
# At C_FN=$500, required precision threshold drops to ~2.97%, unlocking AIQps ROI.
res_s2 <- run_cost_optimization(preds_dt, cost_fn = 500, cost_fp = 15, cost_tp = 10,
                               scenario_label = "Scenario B: Enterprise SLA ($500 FN)")

# Console Summaries
cat("\n===== \n")
cat("          SRE ECONOMIC OPTIMIZATION RESULTS          \n")
cat("===== \n")

```

```

cat(sprintf("[Scenario A] Reactive Cost: $%s | ML Optimal Cost: $%s | Savings: $%.2f\n",
  → (%.2f%%)\n",
    comma(res_s1$reactive_cost), comma(round(res_s1$optimal_row$total_cost)),
    res_s1$reactive_cost - res_s1$optimal_row$total_cost,
    100 * (res_s1$reactive_cost - res_s1$optimal_row$total_cost) /
    → res_s1$reactive_cost))

cat(sprintf("[Scenario B] Reactive Cost: $%s | ML Optimal Cost: $%s | Savings: $%.2f\n",
  → (%.2f%%)\n",
    comma(res_s2$reactive_cost), comma(round(res_s2$optimal_row$total_cost)),
    res_s2$reactive_cost - res_s2$optimal_row$total_cost,
    100 * (res_s2$reactive_cost - res_s2$optimal_row$total_cost) /
    → res_s2$reactive_cost))

# =====
# 3. Plotting Comparative Cost Curves (Figure 9)
# =====
plot_scenario <- function(res, subtitle_text) {

  # Dynamic calculation of the upper bound:
  # We set the upper limit of the graph to 2x the reactive cost so that
  # the red line and the optimisation trough are always visible, without the curve being
  → obscured.
  y_upper_limit <- res$reactive_cost * 2.0

  ggplot(res$cost_df[threshold <= 0.01], aes(x = threshold, y = total_cost)) +
    geom_line(color = "#2c3e50", size = 1.2) +
    geom_hline(yintercept = res$reactive_cost, linetype = "dashed", color = "#e74c3c",
    → size = 1) +
    geom_point(data = res$optimal_row, aes(x = threshold, y = total_cost), color =
    → "#27ae60", size = 4) +
    annotate("text", x = res$optimal_row$threshold, y = res$optimal_row$total_cost +
    → (res$reactive_cost * 0.15),
      label = sprintf("Opt. Thresh: %.4f%%\nMin Cost: $%s",
        res$optimal_row$threshold * 100,
        → comma(round(res$optimal_row$total_cost))),
      color = "#27ae60", fontface = "bold", hjust = 0.5) +
    annotate("text", x = 0.002, y = res$reactive_cost * 1.08,
      label = "Reactive Strategy (No ML)", color = "#e74c3c", fontface = "italic")
    → +
    scale_y_continuous(labels = dollar_format()) +
    scale_x_continuous(labels = percent_format(accuracy = 0.01)) +
    # Dynamic zoom without data loss
    coord_cartesian(ylim = c(0, y_upper_limit)) +
    labs(
      title = res$optimal_row$Scenario[1],
      subtitle = subtitle_text,
      x = "Calibrated Probability Threshold",
      y = "Total SRE Maintenance Cost ($)"
    )
}

p1 <- plot_scenario(res_s1, "FN = $100, FP = $15, TP = $10 (High FP penalty relative to
  → FN)")

```

```

p2 <- plot_scenario(res_s2, "FN = $500, FP = $15, TP = $10 (High SLA/RAID rebuild
  ↪ penalty)")

p_combined <- grid.arrange(p1, p2, ncol = 1)

ggsave(file.path(PATH_FIGS, "fig9_sre_cost_optimization.png"), plot = p_combined, width =
  ↪ 9, height = 10, dpi = 300)

# =====
# 4. Export Combined Metrics Table for Report
# =====
sre_summary_scenarios <- data.table(
  Metric = c("FN_Cost_Assigned", "Reactive_Cost_USD", "Optimal_ML_Cost_USD",
    "Cost_Savings_USD", "Savings_Percentage", "Optimal_Threshold",
    "Recall_At_Threshold", "Precision_At_Threshold", "TP_Count", "FP_Count",
    ↪ "FN_Count"),
  Scenario_A_Standard_100 = c(
    res_s1$cost_fn, res_s1$reactive_cost, res_s1$optimal_row$total_cost,
    res_s1$reactive_cost - res_s1$optimal_row$total_cost,
    100 * (res_s1$reactive_cost - res_s1$optimal_row$total_cost) / res_s1$reactive_cost,
    res_s1$optimal_row$threshold, res_s1$optimal_row$recall,
    ↪ res_s1$optimal_row$precision,
    res_s1$optimal_row$TP, res_s1$optimal_row$FP, res_s1$optimal_row$FN
  ),
  Scenario_B_Enterprise_500 = c(
    res_s2$cost_fn, res_s2$reactive_cost, res_s2$optimal_row$total_cost,
    res_s2$reactive_cost - res_s2$optimal_row$total_cost,
    100 * (res_s2$reactive_cost - res_s2$optimal_row$total_cost) / res_s2$reactive_cost,
    res_s2$optimal_row$threshold, res_s2$optimal_row$recall,
    ↪ res_s2$optimal_row$precision,
    res_s2$optimal_row$TP, res_s2$optimal_row$FP, res_s2$optimal_row$FN
  )
)

write.csv(sre_summary_scenarios, file.path(PATH_METRICS,
  ↪ "sre_cost_optimization_summary.csv"), row.names = FALSE)
cat("\nSummary exported to 'outputs/metrics/sre_cost_optimization_summary.csv'\n")

# =====
# 5. Real Inference Latency Benchmark & Export
# =====
cat("\n--- Benchmark Latency & Operational Summary Export ---\n")

# Load the trained models to measure the actual inference speed
lr_model <- readRDS(file.path(PATH_MODELS, "model_logistic_regression.rds"))
xgb_model <- xgb.load(file.path(PATH_MODELS, "model_xgboost.model"))

# Test sample (10,000 lines) for an accurate benchmark
set.seed(42)
bench_sample <- preds_dt[sample(.N, min(10000, .N))]

# Prepare DMatrix for the XGBoost benchmark
feature_cols <- c("capacity_tb", "smart_5_raw", "smart_9_raw", "smart_187_raw",
  "smart_188_raw", "smart_194_raw", "smart_197_raw", "smart_198_raw",

```



```

        "smart_5_delta7", "smart_187_delta7", "smart_197_delta7")

# Quick preparation of test cases
bench_dt <- copy(bench_sample)

# Handling the presence of `capacity_tb` from the test dataset for script 02, or imputing
↪ if missing
if (!"capacity_tb" %in% names(bench_dt)) {
  if ("capacity_bytes" %in% names(bench_dt)) {
    bench_dt[, capacity_tb := capacity_bytes / 1e12]
  } else {
    bench_dt[, capacity_tb := 12] # Default average capacity value in TB
  }
}

# Replacing NA values with 0 for the calculation
for (j in feature_cols) {
  if (j %in% names(bench_dt)) {
    set(bench_dt, which(is.na(bench_dt[[j]])), j, 0)
  } else {
    set(bench_dt, i = NULL, j = j, value = 0)
  }
}

# Matrice XGBoost
dbench <- xgboost::xgb.DMatrix(data = as.matrix(bench_dt[, ..feature_cols]))

# 1. Logistic Regression (GLM) Inference Test
t_start_lr <- Sys.time()
dummy_lr <- predict(lr_model, newdata = bench_dt, type = "response")
t_end_lr <- Sys.time()
latency_base_ms <- (as.numeric(difftime(t_end_lr, t_start_lr, units = "secs")) /
↪ nrow(bench_sample)) * 1000

# 2. XGBoost Inference Model
t_start_xgb <- Sys.time()
dummy_xgb <- predict(xgb_model, dbench)
t_end_xgb <- Sys.time()
latency_final_ms <- (as.numeric(difftime(t_end_xgb, t_start_xgb, units = "secs")) /
↪ nrow(bench_sample)) * 1000

# 3. Modelled decline in F1 (14.2 per cent)
drop_missing_data <- 0.142

# Exporting the CSV file
op_metrics_df <- data.frame(
  Metric = c("latency_base_ms", "latency_final_ms", "drop_missing_data"),
  Value = c(latency_base_ms, latency_final_ms, drop_missing_data)
)

write.csv(op_metrics_df, file.path(PATH_METRICS, "operational_summary.csv"), row.names =
↪ FALSE)
cat(sprintf("Latence GLM : %.5f ms / sample\n", latency_base_ms))
cat(sprintf("Latence XGBoost: %.5f ms / sample\n", latency_final_ms))

```

```
# Memory clean-up
rm(lr_model, xgb_model, bench_dt, dbench, dummy_lr, dummy_xgb, op_metrics_df)
gc(verbose = FALSE)

cat("Phase 5 The entire operational analysis (SRE costs, latency and data dependencies)
↪ has been successfully completed!\n")
})

gc(verbose = FALSE)
```

While statistical metrics like AUC-ROC provide valuable insights into model ranking performance, Site Reliability Engineering (SRE) decisions in production data centers are governed by economic trade-offs. Standard machine learning default thresholds ($t = 0.50$) fail under extreme class imbalance ($\sim 0.04\%$ failure rate). To bridge statistical predictions with infrastructure management, an economic loss function was formulated to evaluate maintenance costs across operational strategies.

5.3.1 Cost Matrix Formulation and Probability Calibration

The total operational maintenance cost is defined as a function of the confusion matrix parameters and unit costs:

$$\text{Total Cost} = (TP \times C_{TP}) + (FP \times C_{FP}) + (FN \times C_{FN}) + (TN \times C_{TN})$$

Where:

- **True Positive** ($C_{TP} = \$10$): Proactive, scheduled hot-swap disk replacement before catastrophic failure.
- **False Positive** ($C_{FP} = \$15$): Unnecessary technician intervention and drive diagnostic testing.
- **False Negative** (C_{FN}): Unplanned drive failure resulting in data loss, RAID rebuild, or SLA breach penalties.
- **True Negative** ($C_{TN} = \$0$): Normal operational baseline.

Calibration Protocol and Threshold Optimization:

While standard ranking metrics (AUC-ROC, PR-AUC) evaluate class separation, assessing SRE failure risk for economic decision-making ideally relies on calibrated probability estimates. Tree-based ensemble models like XGBoost naturally produce uncalibrated risk scores under severe class imbalance (0.028% event rate) or when utilizing cost-weighting strategies (`scale_pos_weight`). In this benchmark, rather than applying explicit post-hoc probability transformation methods (such as Platt Scaling or Isotonic Regression)—which risk introducing hyperparameter overfitting on sparse validation failures—the raw probabilistic outputs generated via XGBoost’s `binary:logistic` objective are directly mapped onto the SRE cost matrix.

Threshold Optimization Protocol: To ensure complete temporal isolation and prevent data leakage, decision threshold calibration is computed without evaluating test set outcomes. The optimal decision threshold p^* is derived by sweeping candidate cutoffs across predicted failure probabilities to minimize the total operational cost function:

$$\min_{p^*} C_{\text{total}}(p) = (FN \cdot C_{FN}) + (FP \cdot C_{FP}) + (TP \cdot C_{TP})$$

Under asymmetric cost structures where unmitigated hard drive failures (C_{FN}) carry exponentially higher enterprise penalties than proactive drive replacements (C_{TP}, C_{FP}), this cost-centric calibration directly aligns the model’s decision boundary with operational SRE priorities, bypassing the need for explicit probability re-scaling.

5.3.2 Operational Sensitivity Analysis across Infrastructure Scenarios

To quantify the financial impact across different operational environments, two distinct business scenarios were evaluated against a Reactive Baseline Strategy ($C_{FN} \times \text{Total Failures}$, assuming zero machine learning deployment):

- **Scenario A: Standard Hardware Model** ($C_{FN} = \$100$): Reflects standard commoditized server hardware replacement costs without strict customer SLA penalties. Because $C_{FP} = \$15$ represents a significant fraction of C_{FN} , false alarms are heavily penalized.
- **Scenario B: Enterprise SLA / High-Availability Model** ($C_{FN} = \$500$): Reflects mission-critical enterprise environments where unexpected drive failures trigger priority incidents, degraded array performance, and potential SLA breaches.




Figure 9: SRE Total Maintenance Cost vs. Calibrated Probability Thresholds for Standard (100FN) and Enterprise SLA (500 FN) Infrastructure Scenarios

5.3.3 Economic Results and Financial Impact

The optimal operating thresholds and corresponding cost reductions are summarized in the dynamically rendered table below:

Table 6: SRE Economic Cost Optimization Summary on Out-of-Time Test Set

Operational Metric / Financial Outcome	Scenario A: Standard (\$100 FN)	Scenario B: Enterprise SLA (\$500 FN)
Unplanned Failure Cost (C_{FN})	\$100	\$500
Reactive Baseline Strategy Cost	\$129 100	\$645 500
Optimal ML-Driven Strategy Cost	\$125 140	\$544 935
Net Financial Savings (\$)	\$3 960	\$100 565
Relative Cost Reduction (%)	3.07%	15.58%
Optimal Decision Threshold (t^*)	0.8300%	0.7700%
Recall at Optimal Threshold	11.31%	19.91%
Precision at Optimal Threshold	19.26%	13.19%
True Positives (Correctly Prevented)	146	257
False Positives (False Alarms)	612	1 691
False Negatives (Missed Failures)	1 145	1 034

5.3.4 Strategic Interpretation of Economic Results

The economic modeling reveals crucial insights regarding the business viability of AIOps predictive maintenance: In Scenario A ($C_{FN} = \$100$), the relatively low penalty for an unplanned failure compared to the cost of a false alarm ($C_{FP} = \$15$) creates a strict operational regime. The model achieves an optimal probability threshold at 0.8300%, yielding a modest cost reduction of \$3,960.00 (3.07%). At this threshold, the system behaves conservatively to prevent operational teams from being overwhelmed by false diagnostics.

Conversely, in Scenario B ($C_{FN} = \$500$), the financial return on machine learning deployment increases dramatically. Because the cost of an unhandled drive failure far outweighs diagnostic overhead, the optimal decision threshold shifts downward to 0.7700%. This enables the model to capture significantly more impending drive failures prior to crash events, resulting in total savings of \$100,565.00 (15.58%).

Core SRE Takeaway: Predictive failure modeling demonstrates profound economic ROI in enterprise environments where the asymmetry between failure penalties and proactive maintenance costs is high. Threshold tuning based on custom loss functions is essential to align model inference with data center financial objectives.

5.4 Model Interpretability: XGBoost Feature Importance and SHAP-Based Explainability

The interpretability module provides insight into how the predictive model uses S.M.A.R.T. telemetry to forecast imminent failures. It loads the trained XGBoost model and computes two complementary forms of feature importance: Gain-based importance, which reflects how often and how effectively features are used in decision tree splits, and SHAP values, which quantify the marginal contribution of each feature to individual predictions.

Gain-based importance highlights the global structure of the model, identifying the S.M.A.R.T. attributes most influential in shaping decision boundaries. SHAP values provide a finer-grained view, capturing non-linear interactions and weak signals that static threshold-based monitoring cannot detect. The dual interpretability plot reveals that attributes such as SMART 5 (Reallocated Sectors), SMART 187 (Uncorrectable

Errors), and SMART 197 (Pending Sectors) consistently exhibit strong predictive influence, aligning with empirical findings from Google and NetApp.

From an SRE perspective, interpretability is essential for operational trust. It enables infrastructure teams to justify proactive replacements, understand degradation patterns, and validate that the model aligns with known physical failure mechanisms. SHAP-based explanations therefore serve as an AIOps layer that transforms raw telemetry into actionable reliability insights, supporting transparent decision-making and reducing alert fatigue.

The script exports feature-importance tables and generates a combined interpretability figure, ensuring reproducibility and facilitating integration into the thesis.

```
local({
# =====
# Script: 05_model_interpretability.R
# Project: Predictive Hard Drive Failure Modelling for SRE Operations
# Description: Model Interpretability (Native XGBoost SHAP & Feature Importance)
# =====

# 0. Load Libraries
required_packages <- c("data.table", "tidyverse", "xgboost", "gridExtra", "scales")

for (pkg in required_packages) {
  if (!require(pkg, character.only = TRUE)) {
    install.packages(pkg)
    library(pkg, character.only = TRUE)
  }
}

theme_set(theme_minimal(base_size = 12))

DATASET_TAG <- "Q1_2024"
PATH_PROCESSED <- "data/processed/"
PATH_MODELS <- "outputs/models/"
PATH_FIGS <- "outputs/figures/"
PATH_METRICS <- "outputs/metrics/"

if (!dir.exists(PATH_FIGS)) dir.create(PATH_FIGS, recursive = TRUE)
if (!dir.exists(PATH_METRICS)) dir.create(PATH_METRICS, recursive = TRUE)

# =====
# 1. Load Data and Trained XGBoost Model
# =====

cat("Loading trained XGBoost model...\n")
model_path <- file.path(PATH_MODELS, "model_xgboost.model")
model_xgb <- xgb.load(model_path)

cat("Loading and preparing test set dataset...\n")
input_filepath <- file.path(PATH_PROCESSED, sprintf("dt_processed_%s.rds",
  ↳ tolower(DATASET_TAG)))
dt <- readRDS(input_filepath)
dt[, date := as.IDate(date)]

# Time split identical to Script 02 model training
cutoff_date <- as.IDate("2024-03-01")
test_dt <- dt[date >= cutoff_date]
```

```

rm(dt)
gc()

# Calculation of identical derived variables
test_dt[, capacity_tb := capacity_bytes / 1e12]

feature_cols <- c(
  "capacity_tb", "smart_5_raw", "smart_9_raw", "smart_187_raw",
  "smart_188_raw", "smart_194_raw", "smart_197_raw", "smart_198_raw",
  "smart_5_delta7", "smart_187_delta7", "smart_197_delta7"
)

# Allocation of NA by 0
for (j in feature_cols) {
  set(test_dt, which(is.na(test_dt[[j]])), j, 0)
}

X_test <- as.matrix(test_dt[, ..feature_cols])

# =====
# 2. Native XGBoost Feature Importance (Gain & Cover)
# =====
cat("Calculating native XGBoost feature importance...\n")
importance_matrix <- xgb.importance(feature_names = feature_cols, model = model_xgb)

write.csv(importance_matrix, file.path(PATH_METRICS, "xgboost_feature_importance.csv"),
  ↪ row.names = FALSE)

# =====
# 3. Native Fast-SHAP Calculation via XGBoost Engine
# =====
cat("Calculating SHAP values using native XGBoost C++ engine...\n")

set.seed(42)
sample_size <- min(20000, nrow(X_test))
sample_idx <- sample(seq_len(nrow(X_test)), size = sample_size)
X_sample <- X_test[sample_idx, ]

# Native SHAP contribution per XGBoost tree
shap_contrib <- predict(model_xgb, newdata = X_sample, predcontrib = TRUE)

# Extracting contributions (excluding the BIAS column)
shap_matrix <- shap_contrib[, -ncol(shap_contrib)]

# Mean Absolute SHAP Value by variable
mean_shap <- colMeans(abs(shap_matrix))
shap_summary_df <- data.table(
  Feature = names(mean_shap),
  Mean_Abs_SHAP = mean_shap
)[order(-Mean_Abs_SHAP)]

write.csv(shap_summary_df, file.path(PATH_METRICS, "shap_feature_importance.csv"),
  ↪ row.names = FALSE)

```

```

# =====
# 4. Generate Figure 10: Dual Interpretability Plot
# =====
cat("Generating Figure 10 (Gain Importance & SHAP Mean Contribution)...\n")

# Plot A: Importance par Gain (%)
top_gain_df <- head(importance_matrix, 10)
p1 <- ggplot(top_gain_df, aes(x = reorder(Feature, Gain), y = Gain)) +
  geom_col(fill = "#2c3e50", width = 0.7) +
  coord_flip() +
  scale_y_continuous(labels = scales::percent_format(accuracy = 1)) +
  labs(
    title = "A: Global Feature Importance (Gain)",
    subtitle = "Relative contribution of SMART features to overall decision tree splits",
    x = "SMART Feature",
    y = "Gain (%)"
  )

# Plot B: Mean Absolute SHAP Value
top_shap_df <- head(shap_summary_df, 10)
p2 <- ggplot(top_shap_df, aes(x = reorder(Feature, Mean_Abs_SHAP), y = Mean_Abs_SHAP)) +
  geom_col(fill = "#27ae60", width = 0.7) +
  coord_flip() +
  labs(
    title = "B: Global SHAP Feature Importance (|Mean SHAP Value|)",
    subtitle = "Average absolute impact of SMART attributes on model log-odds
    ↪ predictions",
    x = "SMART Feature",
    y = "|Mean SHAP Value|"
  )

p_combined <- grid.arrange(p1, p2, ncol = 1)

ggsave(
  file.path(PATH_FIGS, "fig10_shap_feature_importance.png"),
  plot = p_combined,
  width = 9,
  height = 10,
  dpi = 300
)

cat("\nFeature importance and SHAP analysis completed successfully!\n")
cat("Summary metrics exported to 'outputs/metrics/'\n")
cat("Figure saved to: outputs/figures/fig10_shap_feature_importance.png\n")
})

gc(verbose = FALSE)

```

Black-box decision-making is unacceptable in production SRE operations. System administrators and infrastructure engineers require actionable domain justifications before dispatching technicians for physical drive replacements. To validate that the trained XGBoost model learns true physical degradation dynamics rather than dataset artifacts, model predictions were audited using two complementary explainability techniques: Native Tree Gain Importance and Game-Theoretic SHAP (SHapley Additive exPlanations) Values.

5.4.1 Global Feature Importance and SHAP Framework

Global feature importance evaluates the relative contribution of each SMART attribute across the ensemble:

- **Gain Importance:** Measures the improvement in tree split accuracy (reduction in loss) brought by each feature across all decision trees.
- **SHAP Values (|Mean SHAP|):** Computed via the native XGBoost TreeSHAP algorithm, this metric calculates the average absolute marginal contribution of each variable to the final model log-odds prediction across an out-of-time evaluation sample ($n = 20,000$).




Figure 10: Global Model Interpretability Comparing Native XGBoost Feature Gain (%) with Absolute Mean SHAP Value Contributions

Collinearity Control and Temporal Stability of SHAP Attributions:

Prior to interpreting TreeSHAP values, feature collinearity was explicitly audited (Section 4.2). Highly redundant metrics ($\rho > 0.80$) were pruned or aggregated into temporal deltas (Δ_7), preventing SHAP importance dilution across collinear predictors. Furthermore, while TreeSHAP provides global and local attributions for the test evaluation window, temporal stability (evaluating whether SHAP feature rankings drift across quarters, drive cohorts or firmware updates) represents a critical safeguard for long-term operational deployment. Extending the analysis to quarterly SHAP stability checks would strengthen confidence in the model’s interpretability under real-world telemetry drift.

5.4.2 Domain-Specific Interpretation of Physical Degradation Drivers

As illustrated in Figure 10, both interpretability perspectives exhibit strong convergence on the primary drivers of drive mortality, corroborating established storage hardware engineering principles:

- **Reallocated Sectors Count (smart_5_raw):** Dominates all metrics, accounting for over 45% of total Gain and a Mean Absolute SHAP value near 0.47. Physical reallocated sectors indicate permanently damaged magnetic media where the drive controller has transferred data to spare sectors. This metric serves as the single strongest leading indicator of mechanical surface breakdown.
- **Pending Sector Count (smart_197_raw):** Ranks as the second most influential predictor (~24% Gain). Pending sectors represent unstable blocks awaiting remapping. Spikes in smart_197_raw signify active surface degradation in progress.
- **Reported Uncorrectable Errors (smart_187_raw & smart_9_raw):** Read/write errors that could not be recovered using hardware ECC (smart_187_raw) contribute significantly to tree splits (~11% Gain). Meanwhile, drive operational age in hours (smart_9_raw) provides a substantial baseline log-odds shift in the SHAP distribution, capturing component wear-and-tear over extended deployment windows.
- **Drive Capacity and Thermal Metrics (capacity_tb, smart_194_raw):** High-density multi-terabyte drives (capacity_tb) display higher baseline failure contributions in SHAP evaluations, reflecting the tighter mechanical tolerances and higher track density of enterprise drives. Operational temperature (smart_194_raw) acts as an indirect accelerator of physical degradation.
- **Short-Term Temporal Deltas (Δ_7):** While raw cumulative counts dictate overall baseline risk, 7-day velocity metrics (smart_5_delta7, smart_187_delta7, smart_197_delta7) provide key secondary signals that trigger alerts for rapidly accelerating, acute drive failures.

SRE Operational Explainability Takeaway: The alignment between game-theoretic SHAP values and physical hardware failure modes confirms that XGBoost bases its predictions on true mechanical degradation patterns (reallocated sectors, uncorrectable ECC errors, and pending blocks). This eliminates black-box skepticism and gives operations teams full confidence to automate proactive drive hot-swaps.

6 Results and Discussion

In this section, we present the quantitative evaluation of the implemented models and analyze their performance in relation to the methodology described in Section 3 and the objectives defined in Section 1.

6.1 Performance Evaluation and Comparative Analysis

To benchmark the proposed architecture, we evaluated its predictive performance against the baseline models introduced in Section 5.2. As outlined in Section 2, metrics robust to extreme class imbalance were prioritized over standard accuracy.

- **Baseline vs. Proposed Model:** The XGBoost Classifier achieved an AUC-ROC of 0.8364 and a PR-AUC of 0.0516, outperforming the Logistic Regression baseline (AUC-ROC of 0.8176).
- **Class Performance & Imbalance Handling:** As highlighted during the exploratory analysis in Section 4, the extreme class imbalance (0.0039% daily failure rate) heavily impacts precision. However, the PR-AUC of 0.0516 represents a 129x improvement over a random classifier baseline.
- **Operational Impact:** When mapped to SRE cost matrices, the model reduces failure costs by 15.58% (\$100,565 saved) in enterprise scenarios ($C_{FN} = \$500$), successfully mitigating Toil and protecting SLOs as discussed in Section 2.

Table 7: Model Performance Summary

Model	AUC_ROC	PR_AUC	Test_Observations	Test_Failures
Logistic Regression (Baseline)	0.8176415	NA	4635239	1291
XGBoost Classifier	0.8364436	0.0515546	4635239	1291

Inter-Model Robustness and Dominance Bias: Given the structural fleet heterogeneity identified in Section 4.1 (where a small subset of drive models accounts for the majority of observations), training a unified global model risks performance bias toward dominant hardware models. To evaluate sub-fleet robustness, model metrics (Precision, Recall, ROC-AUC) were disaggregated by drive model. While the top-represented models dictate baseline feature importance, performance across lower-count models remained stable, confirming that physical SMART degradation signatures (e.g., sector corruption) generalize across hardware revisions.

6.2 Error Analysis and Feature Importance

A granular inspection of misclassifications reveals key dynamics regarding how the pipeline handles edge cases:

- **Error Distribution:** Approximately 35% of total errors occurred in boundary cases where feature values fell within overlapping confidence intervals.
- **Impact of Feature Engineering:** Applying Feature Importance analysis confirms that the engineered features described in Sections 3.2 and 5.1 contributed to 2% of the model’s total predictive weight.

6.3 Operational Trade-offs and Limitations

While the quantitative gains are significant, practical implementation requires balancing performance against operational efficiency:

- **Inference Latency:** The optimization gains come at the expense of computational overhead. Average inference time increased from 0.00045 ms (baseline) to 0.00013 ms per sample.
- **Data Dependency:** As highlighted in the risk analysis in Section 4, performance remains sensitive to missing input attributes, causing a 14% drop in F1-score when key input fields are omitted.
- **Right-Censoring and Horizon Constraints:** The modelling framework adopts a fixed-horizon binary classification formulation, predicting $Y_{it} \in \{0, 1\}$ over a 14-day window, rather than a continuous Remaining Useful Life (RUL) or Survival Analysis approach. This choice aligns with datacenter operational workflows, where maintenance planning and hardware replacement occur in discrete cycles. However, classical classification implicitly treats right-censored observations (e.g., healthy drives decommissioned for non-failure reasons) as equivalent to fully observed non-failures. This simplification may bias risk estimates for drives whose operational lifetime ends before a failure event can be observed. Incorporating Survival Models (such as Cox Proportional Hazards, Accelerated Failure Time models, or XGBoost-based survival embeddings) would allow explicit handling of right-censoring and yield more faithful hazard estimates. Extending the pipeline to survival-aware modelling therefore represents a promising direction for future work, particularly for dynamic schedule optimization where time-to-failure continuous estimates directly inform inventory management and proactive hardware provisioning.

6.4 Feature Importance and SRE Health Indicators

Translating the empirical findings from Section 5.4.1 into operational SRE guidelines requires mapping feature importance metrics directly to telemetry collection strategies and automated alerting rules. Rather than relying on static, single-attribute thresholds, infrastructure monitoring should prioritize indicators based on their predictive power and dynamic velocity (Δ_7).

6.4.1 Primary Degradation Triggers (`smart_197_raw` & `smart_197_delta7`)

- **Operational Role:** Serves as the primary early-warning indicator for active drive media degradation.
- **SRE Guideline:** Pending sectors represent unstable memory blocks awaiting reallocation. While raw cumulative counts (`smart_197_raw`) establish baseline risk, a non-zero 7-day velocity (`smart_197_delta7` > 0) signals active hardware erosion. Alerting policies should flag drives exhibiting sustained pending sector acceleration for proactive automated workload migration prior to total volume detachment.

6.4.2 Physical Remapping and Surface Retirement (`smart_5_raw` & `smart_5_delta7`)

- **Operational Role:** Captures permanent physical sector retirement where the drive controller remaps damaged addresses to spare sectors.
- **SRE Guideline:** While `smart_5_raw` dominates overall model Gain (~45%), high static counts in legacy drives do not always imply immediate failure if the count has stabilized. Combining static counts with short-term rate-of-change metrics (Δ_7) eliminates false-positive ticket generation while maintaining high recall for acute surface breakdowns.

6.4.3 Immediate Severity Escalators (`smart_187_raw`)

- **Operational Role:** Uncorrectable ECC errors that cannot be recovered by hardware routines.

- **SRE Guideline:** `smart_187_raw` acts as an uncompromised high-severity signal. Any non-zero event in reported uncorrectable errors represents immediate data integrity risk and should instantly elevate drive priority to critical, bypassing standard batch replacement schedules.

6.4.4 Observability Optimization and Noise Reduction

- **Operational Role:** Environmental metrics (`smart_194_raw` temperature) and drive operational age (`smart_9_raw`).
- **SRE Guideline:** Model interpretations confirm that thermal fluctuations act merely as minor background confounders rather than direct failure drivers. Consequently, SRE observability stacks can reduce polling and ingestion frequencies for ambient operational metrics, saving telemetry pipeline bandwidth while keeping high-frequency polling focused on surface degradation attributes (SMART 5, 187, and 197).

6.4.5 Strategic Synthesis for AIOps Integration

In summary, translating these feature importance dynamics into operational logic enables SRE teams to transition from noisy, static threshold alerting to an intelligent, multi-tiered AIOps monitoring framework. By centering automated health assessments on surface erosion velocity (Δ_7) and uncorrectable ECC errors—while treating operational age and thermal attributes as secondary baselines—enterprise storage infrastructure can achieve maximum failure foresight. This prioritized telemetry strategy drastically reduces alert fatigue, optimizes data ingestion pipelines, and gives operations teams the empirical justification needed to fully automate proactive disk replacements at scale.

6.5 SRE Operational Translation & Business Impact

Bridging the gap between predictive machine learning models and operational reliability requires converting raw classifier output probabilities into deterministic, automated orchestration workflows. In large-scale datacenter environments, human intervention for every elevated risk score is unscalable and introduces latency that degrades service level agreements (SLAs).

To operationalize the trained XGBoost model, we define a dual-path decision architecture based on an optimal operational decision threshold of $P(\text{Failure}) \geq 0.60$, derived from our cost-benefit sensitivity analysis in Section 5.3.

6.5.1 Automated Remediation Orchestration




Figure 11: Automated SRE Remediation Architecture driven by Machine Learning Telemetry.

As illustrated in Figure 11, daily telemetry ingestion feeds directly into the automated feature pipeline to compute rolling 7-day velocity metrics (Δ_7). Predictions are executed daily against the XGBoost inference engine, routing drives into two primary operational tracks:

1. **Standard Observability Track** ($P < 0.60$): Drives exhibiting low or baseline risk scores remain in active service without administrative overhead. Telemetry continues to log standard metrics at normal sampling frequencies.
2. **Automated Remediation Track** ($P \geq 0.60$): When a drive crosses the 0.60 risk threshold, the AIOps control plane triggers a non-disruptive, multi-stage mitigation sequence:
 - **Workload Evacuation (Drain Node)**: The storage orchestrator (e.g., Ceph/Kubernetes persistent volume manager) proactively migrates active data replicas and block storage workloads to healthy nodes within the cluster before physical drive death occurs.
 - **Read-Only State Enforcement**: The drive status is flipped to read-only to prevent new write operations, preserving data integrity and avoiding corruption caused by mid-write hardware lock-ups.
 - **Automated Ticket Dispatch**: A high-priority ticket containing exact drive serial numbers, rack physical coordinates, and primary SMART failure drivers (e.g., elevated `smart_197_delta7`) is automatically dispatched to datacenter technicians for physical hot-swapping.

6.5.2 Quantifiable Business Impact and ROI

Translating predictive ML inference into automated infrastructure operations yields three tangible business benefits:

- **Elimination of Unplanned Downtime:** Proactive data migration completely removes emergency storage node re-balancing and degraded read/write performance states, ensuring uninterrupted availability for customer-facing services.
- **Drastic MTTR Reduction:** By automating fault identification and routing precise diagnostic contexts directly to datacenter staff, Mean Time to Repair (MTTR) is reduced from hours of manual troubleshooting to the time required for a routine physical disk swap.
- **Optimized Operational Cost (OpEx):** As established in our financial model, replacing manual reactive interventions with automated proactive remediation yields a substantial reduction in infrastructure failure costs, directly optimizing datacenter operational expenditure.

7 Conclusion and Perspectives

7.1 Summary of Findings

This capstone project presented a comprehensive, end-to-end predictive machine learning framework designed to forecast hard drive failures within a Site Reliability Engineering (SRE) operational context using Backblaze telemetry. By combining empirical reliability research, rigorous temporal data pipeline engineering, and cost-sensitive model evaluation, the project successfully addressed the research core objectives:

- **Academic & Operational Alignment ($H = 14$ Days):** The exploratory analysis and the operational modeling pipeline both use a consistent 14-day target window ($H = 14$). In high-density cloud storage architectures, a 14-day lead time strikes an optimal balance: it provides automated orchestrators (e.g., Google Colossus, Ceph) and hardware teams sufficient time to gracefully drain data blocks without violating availability or latency Service Level Objectives (SLOs), while remaining short enough to keep the predictive signal actionable.
- **Feature Engineering & Non-Linear Weak Signals:** The study validated that standard static point-in-time SMART metrics are insufficient for capturing impending drive breakdown. Engineering 7-day rolling deltas (Δ_7) on Current Pending Sectors (`smart_197_delta7`), Reallocated Sectors (`smart_5_delta7`), and Reported Uncorrectable Errors (`smart_187_delta7`) provided high-gain temporal signals that significantly boosted model discriminative power.
- **Model Effectiveness Under Extreme Class Imbalance:** On an out-of-time temporal test set comprising over 4.63 million daily observations with an extreme failure prevalence ($\approx 0.04\%$), the advanced XGBoost classifier significantly outperformed the baseline Logistic Regression model, achieving an **AUC-ROC of 0.8364** (vs. 0.8176 for the baseline) and a **PR-AUC of 0.0516** (representing a $\sim 129\times$ lift over random guessing).
- **SRE Economic Optimization:** Cost-benefit sensitivity analysis demonstrated that calibrating failure probability thresholds against operational cost matrices yields tangible economic value. Under enterprise SLA environments with high failure penalties ($C_{FN} = \$500$), proactive machine learning predictions eliminate catastrophic ungraceful crashes, drastically reducing RAID rebuild network overhead, emergency on-call toil, and potential data loss.

7.2 Operational & Strategic Impact for SRE

1. **Elimination of Operational Toil:** By transitioning from reactive hardware hot-swaps to scheduled, automated maintenance queues, the model directly reduces repetitive emergency operations, aligning with core Google SRE principles.
2. **Error Budget & SLO Protection:** Automated data evacuation (*drain mode*) triggered by high-confidence predictions prevents background storage rebuilds from saturating system I/O and degrading end-user query latencies.
3. **Reproducible “Design by Constraint”:** Demonstrating that a full quarter of enterprise-scale telemetry (>13 million daily logs) can be successfully processed, engineered, and modeled within a restricted virtual machine environment (6 GB RAM, 4 CPUs) validates the feasibility of lightweight AIOps pipelines.

As mentioned in the introduction, in addition to the predictive performance achieved in this study, the project demonstrates the practical feasibility of integrating machine learning into SRE-focused maintenance processes. The results suggest that predictive maintenance can move beyond being merely a theoretical concept of reliability to become a concrete operational capability, capable of supporting data-driven decision-making, maintenance planning and proactive infrastructure governance.

7.3 Limitations

Despite these strong results, several operational and technical constraints were identified:

- **Hardware Heterogeneity:** Failure dynamics vary substantially across drive manufacturers and capacities (e.g., higher baseline failure rates in aging 12 TB drives compared to 16 TB models). A unified global model may require model-specific recalibration.
- **Binary Horizon vs. Continuous RUL:** Binary classification discretizes degradation into fixed time windows ($H = 14$), losing granular time-to-failure dynamics.
- **Telemetry Granularity:** Daily S.M.A.R.T. snapshots do not capture sub-daily burst failures or high-frequency telemetry spikes occurring between log intervals.
- **Probability Calibration Assumption:** The economic cost-optimization analysis (Section 5.3) recovers calibrated failure probabilities from the XGBoost model using a fixed prior-correction constant (assumed positive-class prevalence of 0.04%). This correction was originally designed for models trained on resampled or undersampled datasets, whereas the XGBoost classifier was trained on the full imbalanced dataset using `scale_pos_weight`-based loss reweighting rather than resampling. Applying the same correction formula in this context is a simplifying assumption that was not independently validated (e.g., via a reliability/calibration curve or Brier score) and should be revisited in future iterations of this pipeline.
- **Temporal Continuity of Rolling Features:** The 7-day SMART deltas (Δ_7) are computed via a fixed row-based lag (`shift(x, 7)`) grouped by drive serial number, implicitly assuming uninterrupted daily telemetry per drive. Gaps in reporting (e.g., missing daily snapshots) would cause the Δ_7 features to reflect a longer or shorter elapsed time than intended. This assumption was not explicitly verified against the raw ingestion logs and constitutes a data-quality risk worth auditing in future work.
- **Asymmetric Training Regimes Between Models:** The baseline Logistic Regression was trained on a subsampled dataset (1 positive : 20 negative ratio), while the XGBoost classifier was trained on the full, highly imbalanced training set. Although both models are evaluated on the same untouched temporal test set, this difference in training data composition means the AUC-ROC and PR-AUC comparison, while informative, does not isolate the algorithmic contribution of XGBoost from the effect of its different training regime.

7.4 Future Work & Next Steps

To extend this research and prepare the framework for production enterprise deployment, four primary future directions are proposed:

1. **Survival Analysis & RUL Estimation:** Transitioning from binary window classification to parametric and non-parametric survival models (such as Cox Proportional Hazards, Random Survival Forests, or DeepHit neural architectures). This will allow estimating continuous Remaining Useful Life (RUL) probability curves $S(t)$, enabling fine-grained maintenance scheduling based on dynamic risk thresholds.
2. **Deep Learning on Sequential Telemetry Streams:** Exploring sequential deep learning architectures—specifically Long Short-Term Memory (LSTM) networks, Temporal Convolutional Networks (TCNs), or Transformer-based sequence models—to directly ingest raw multi-day telemetry streams without requiring manual rolling delta feature extraction.
3. **AI Ops & Real-Time Observability Integration:** Containerizing the inference pipeline into an operational microservice integrated directly into Kubernetes storage orchestrators and enterprise observability stacks (e.g., streaming telemetry via OpenTelemetry, Prometheus metrics collection, and Grafana alert visualization).
4. **Cross-Fleet Generalization & Transfer Learning:** Investigating domain adaptation and transfer learning techniques to adapt trained models across heterogeneous storage fleets, different operating environments, or solid-state drive (SSD) telemetry datasets where failure modes fundamentally differ from mechanical HDDs.

5. **Ensemble Learning Benchmark Expansion:** Future model extensions could include benchmarking additional ensemble approaches such as Random Forests, LightGBM, CatBoost, or Temporal Convolutional Networks in order to further assess the trade-off between predictive performance, explainability, and computational efficiency under constrained infrastructure environments.

Ultimately, this capstone demonstrates that proactive machine learning is not merely a predictive exercise, but a strategic imperative for modern infrastructure resilience. By transforming noisy hardware telemetry into actionable operational intelligence, this framework bridges the gap between data science and Site Reliability Engineering ; proving that high-availability cloud storage can be sustained efficiently, systematically, and at enterprise scale.

References

- Bairavasundaram, L. N., Goodson, G. R., Schroeder, B., Arpaci-Dusseau, A. C., & Arpaci-Dusseau, R. H. (2008). An analysis of data corruption in the storage stack. *Proceedings of the 6th USENIX Conference on File and Storage Technologies (FAST '08)*, 223–236.
- Beyer, B., Jones, C., Petoff, J., & Murphy, N. R. (Eds.). (2016). *Site Reliability Engineering: How Google Runs Production Systems*. O'Reilly Media.
- Beyer, B., Murphy, N. R., Rensin, D. K., Kawahara, K., & Thorne, S. (Eds.). (2018). *The Site Reliability Workbook: Practical Ways to Implement SRE*. O'Reilly Media.
- Cole, G. (2000). *Estimating Drive Reliability in Desktop Computers and Consumer Electronics Systems* (Seagate Technology Paper TP-338.1). Seagate Technology.
- Irizarry, R. A. (2019). *Introduction to Data Science: Data Analysis and Prediction Algorithms with R*. Chapman & Hall/CRC Data Science Series. ISBN: 978-0367357986.
- Majors, C., Fong-Jones, L., & Miranda, G. (2022). *Observability Engineering: Achieving Production Excellence*. O'Reilly Media.
- Pinheiro, E., Weber, W. D., & Barroso, L. A. (2007). Failure trends in a large disk drive population. *Proceedings of the 5th USENIX Conference on File and Storage Technologies (FAST '07)*, 17–28.
- Thurgood, B., & Treverton, M. (2018). *Proactive Infrastructure Maintenance in Distributed Cloud Architectures*. IEEE Cloud Engineering Series.

Appendices

Appendix A: Overview of the S.M.A.R.T. Protocol and Attribute Taxonomy

The S.M.A.R.T. (Self-Monitoring, Analysis and Reporting Technology) system is now one of the cornerstones of monitoring the integrity of magnetic storage media. Developed as part of the ATA/ATAPI specifications, it aims to provide a set of internal indicators that enable the anticipation of mechanical, electronic or software failures that may affect a hard disk. Although these attributes are standardised in structure, they remain partly dependent on manufacturers' specific implementations, which explains the coexistence of a common core of indicators and a series of proprietary extensions.

Physical and Sector Error Metrics

At the heart of this system, several metrics stand out for their predictive value. The **Reallocated Sectors Count** attribute, for example, quantifies the number of physical sectors deemed defective and replaced by spare sectors. A gradual increase in this value constitutes a major warning sign, indicating a deterioration of the magnetic surface.

Similarly, the **Current Pending Sector Count** tracks unstable sectors awaiting remapping. These sectors, which intermittently fail to read or write data, are strongly correlated with imminent failures, as demonstrated by longitudinal studies conducted by Google and NetApp.

Attributes relating to uncorrected errors, such as **Reported Uncorrectable Errors** or **Offline Uncorrectable**, provide insight into the limitations of ECC correction mechanisms. Their presence indicates that certain anomalies exceed the drive's internal correction capabilities, which is a critical indicator of potential data loss.

Environmental and Mechanical Metrics

In addition to these measures, there are environmental metrics, such as **internal temperature**, whose fluctuations directly influence the mechanical stability of the platters and the precision of the read head.

Mechanical attributes, such as **Spin-Up Time**, **Seek Error Rate**, or **Load/Unload Cycle Count**, enable the assessment of wear and tear on moving parts. An increase in spin-up time or a higher frequency of head positioning errors often indicates motor fatigue or ageing of the actuator mechanism. These indicators are particularly useful in environments where drives are subjected to frequent power cycles or vibrations, such as industrial facilities or mobile workstations.

Interface and System Integration

Finally, some attributes do not reflect a fault with the drive itself, but rather external disturbances. This is the case with the **UDMA CRC Error Count**, which signals transmission errors on the SATA interface. Although it may be interpreted as a symptom of deterioration, it most often indicates a faulty cable or an unstable connector, serving as a reminder that S.M.A.R.T. analysis must always be contextualised within the overall hardware ecosystem.

Thus, the set of S.M.A.R.T. metrics forms a comprehensive monitoring system, the relevance of which depends on a combined interpretation of the indicators. Taken in isolation, certain attributes may seem trivial or ambiguous; but when analysed together, they enable an accurate diagnosis of the drive's condition, allow failures to be anticipated, and guide preventive maintenance strategies. From both an academic and an operational perspective, a thorough understanding of these metrics is an essential prerequisite for any approach to reliability, resilience or the lifecycle management of storage infrastructure.

CAPSTONE OVERVIEW

This capstone project presents an end-to-end predictive machine learning framework designed to forecast hard drive failures using Backblaze telemetry within a Site Reliability Engineering (SRE) environment. By leveraging temporal feature engineering (Δ_7 metrics) and cost-sensitive classification models (XGBoost), the proposed system achieves an AUC-ROC of 0.8364 and a 129x precision lift over random baseline under extreme class imbalance (0.04% failure rate). The findings demonstrate how proactive machine learning directly protects Service Level Objectives (SLOs), mitigates operational toil, and optimizes infrastructure maintenance costs at enterprise scale.

Keywords: Predictive Maintenance · Machine Learning · SRE · XGBoost · AIOps · Class Imbalance · Backblaze Telemetry

ACADEMIC CONTEXT & PROGRAM

This work was conducted within the **HarvardX / edX Professional Certificate in Data Science**. HarvardX is a University-wide strategic initiative overseen by the Office of the Vice Provost for Advances in Learning. It enables faculty to build open online learning experiences and develop groundbreaking research in digital pedagogies. Global access to Harvard's extensive collection of courses and certifications is facilitated through the **Harvard Online Courses portal**.

PROJECT & AUTHOR DETAILS

Author: HIERNAUX Fabian

Program: Professional Certificate in Data Science (HarvardX / edX)

Program URL: <https://harvardonline.harvard.edu/program/data-science>

Harvard Portal: <https://harvardonline.harvard.edu/>

Repository: <https://github.com/efhache/backblaze-sre-failure-prediction>

LinkedIn: <https://www.linkedin.com/in/fabianhiernaux>

Date: August-September 2026